

# The Graphics Pipeline and OpenGL IV:

Stereo Rendering, Depth of Field Rendering, Multi-pass Rendering

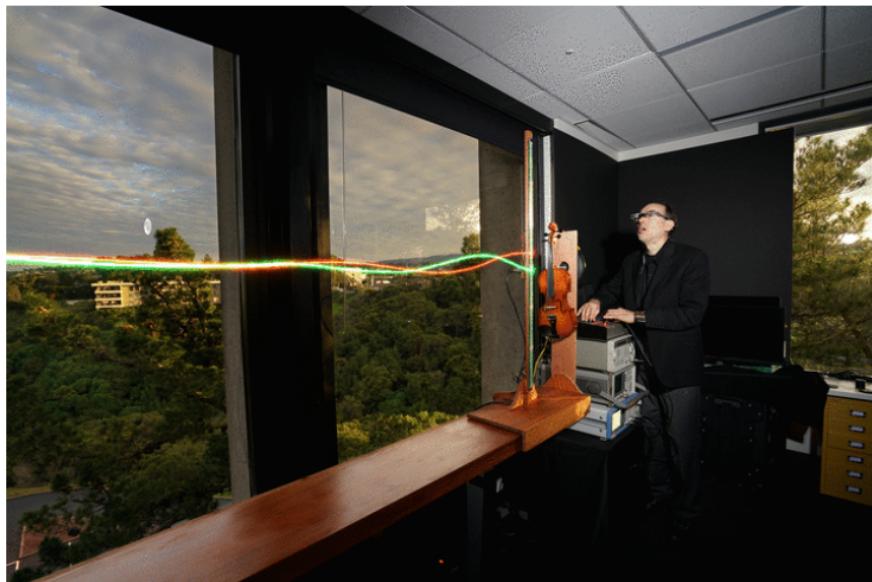


Gordon Wetzstein  
Stanford University

EE 267 Virtual Reality  
Lecture 6

[stanford.edu/class/ee267/](http://stanford.edu/class/ee267/)

# Talk by Steve Mann today 4:30pm, Packard 101



# Lecture Overview

- overview of glasses-based stereo
- stereo rendering with OpenGL
  - projection matrix
  - view matrix
- offscreen frame buffers and multi-render passes
- anaglyph stereo rendering with GLSL
- depth of field rendering

# Glasses-based Stereo



1. Anaglyph



2. Polarization



3. Shutter Glasses

4. Chromatic Filters (Dolby)



# Glasses-based Stereo



2. Polarization

# Glasses-based Stereo

- passive glasses
- active LC element on projector or interlaced rows/columns on monitor (resolution loss)



## 2. Polarization

- e.g. RealD – most 3D cinemas use this
- circular polarization to allow for head roll
- inexpensive glasses, little crosstalk
- need polarization-preserving screen!



# Glasses-based Stereo



3. Shutter Glasses

# Glasses-based Stereo

- active glasses, temporally-multiplexed display
- e.g. StereoGraphics
- somewhat expensive glasses, little crosstalk
- need fast display (at least 120 Hz)
- sync monitor update with glasses



3. Shutter Glasses

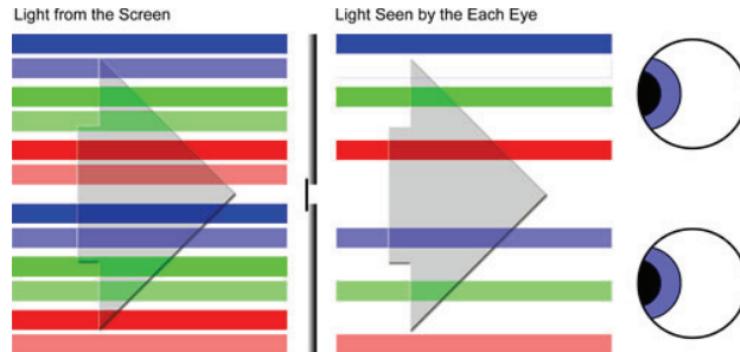
# Glasses-based Stereo

4. Chromatic Filters (Dolby)



# Glasses-based Stereo

- passive glasses, usually two projectors with passive color filters
- somewhat expensive glasses (not as widespread in cinemas)
- full color!



4. Chromatic Filters (Dolby)



# Glasses-based Stereo



1. Anaglyph

# Glasses-based Stereo



## 1. Anaglyph

- passive, inexpensive glasses (least expensive overall)
- no modifications to display necessary – just render stereo images in different colors
- cannot reproduce correct colors! but not as bad as it sounds



Put on Your 3D Glasses Now!



# Anaglyph Stereo - Monochrome

- render L & R images, convert to grayscale
- merge into red-cyan anaglyph by assigning  $I(r)=L$ ,  $I(g,b)=R$  ( $I$  is anaglyph)



from movie "Bick Buck Bunny"



# Anaglyph Stereo – Full Color

- render L & R images, do not convert to grayscale
- merge into red-cyan anaglyph by assigning  $I(r)=L(r)$ ,  $I(g,b)=R(g,b)$  ( $I$  is anaglyph)



from movie "Bick Buck Bunny"





# Anaglyph Stereo - Dubois

- paper: Eric Dubois, ICASSP 2001, “A Projection Method to Generate Anaglyph Stereo Images”
- optimize color management in CIE XYZ space
- requires spectral transmission of glasses & spectral emission curves of display primaries
- great course project - see last year’s projects ...

Open Source Movie: Big Buck Bunny

Rendered with Blender (Open Source 3D Modeling Program)

<http://bbb3d.renderfarming.net/download.html>

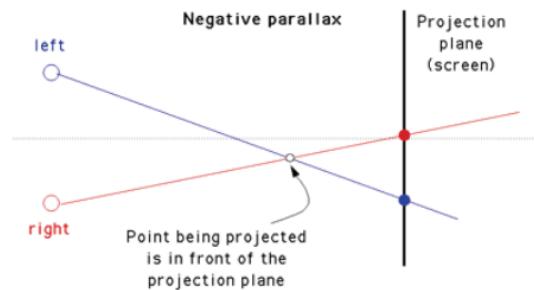
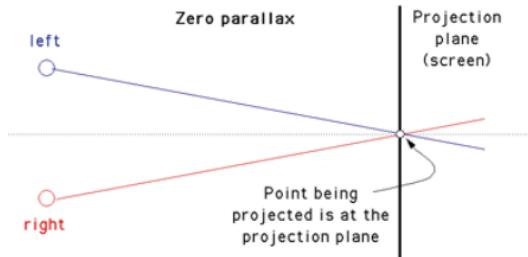
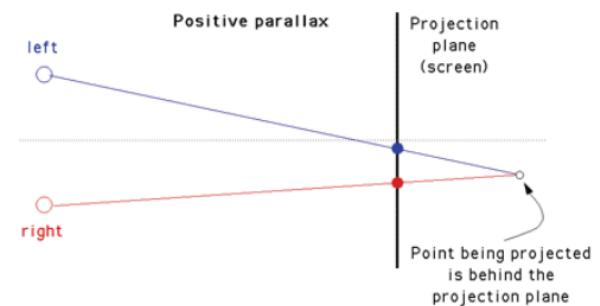


# Anaglyph Comparison

... show video clips ...

# Parallax

- parallax is the relative distance of a 3D point projected into the 2 stereo images



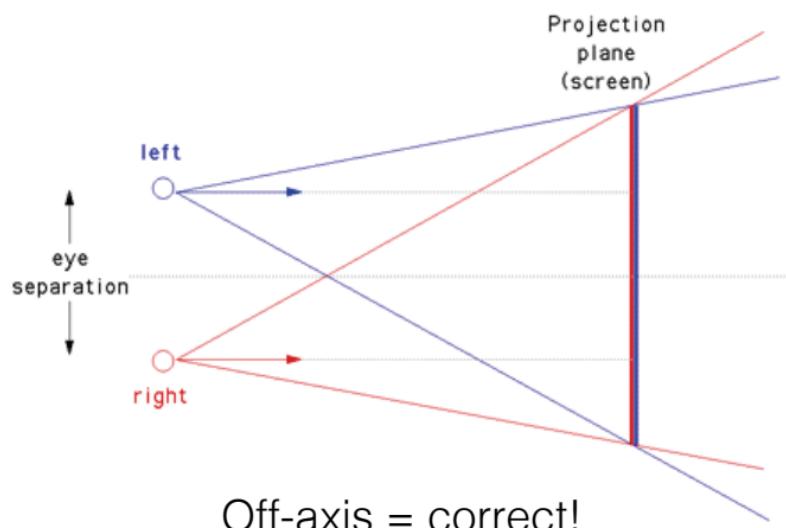
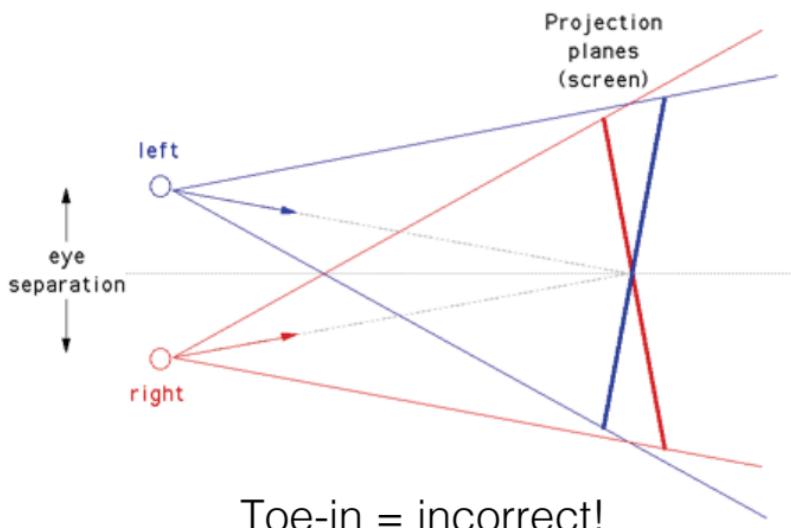
case 1

case 2

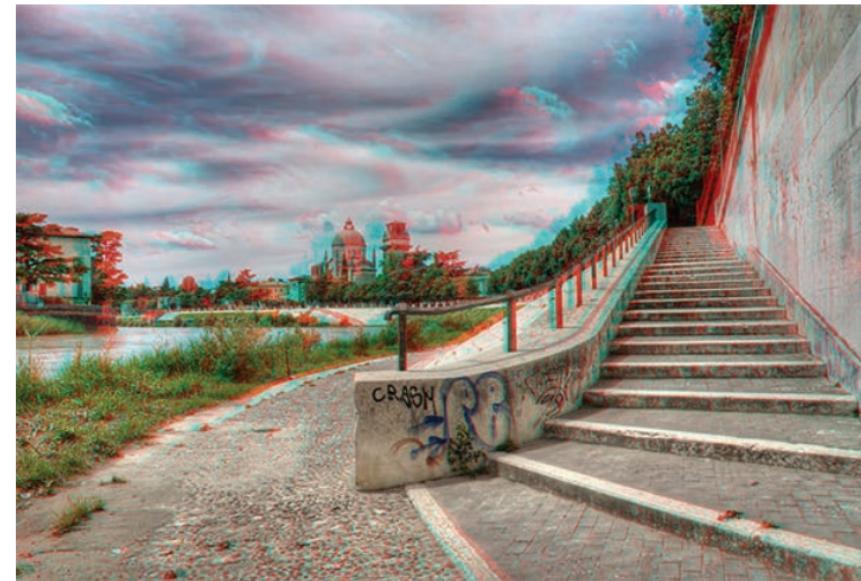
case 3

# Parallax

- visual system only uses horizontal parallax, no vertical parallax!
- naïve toe-in method creates vertical parallax → visual discomfort



# Parallax – well done



# Parallax – well done



1862

"Tending wounded Union soldiers at  
Savage's Station, Virginia, during the  
Peninsular Campaign",  
Library of Congress Prints and  
Photographs Division



# Parallax – not well done

- vertical parallax really doesn't work!





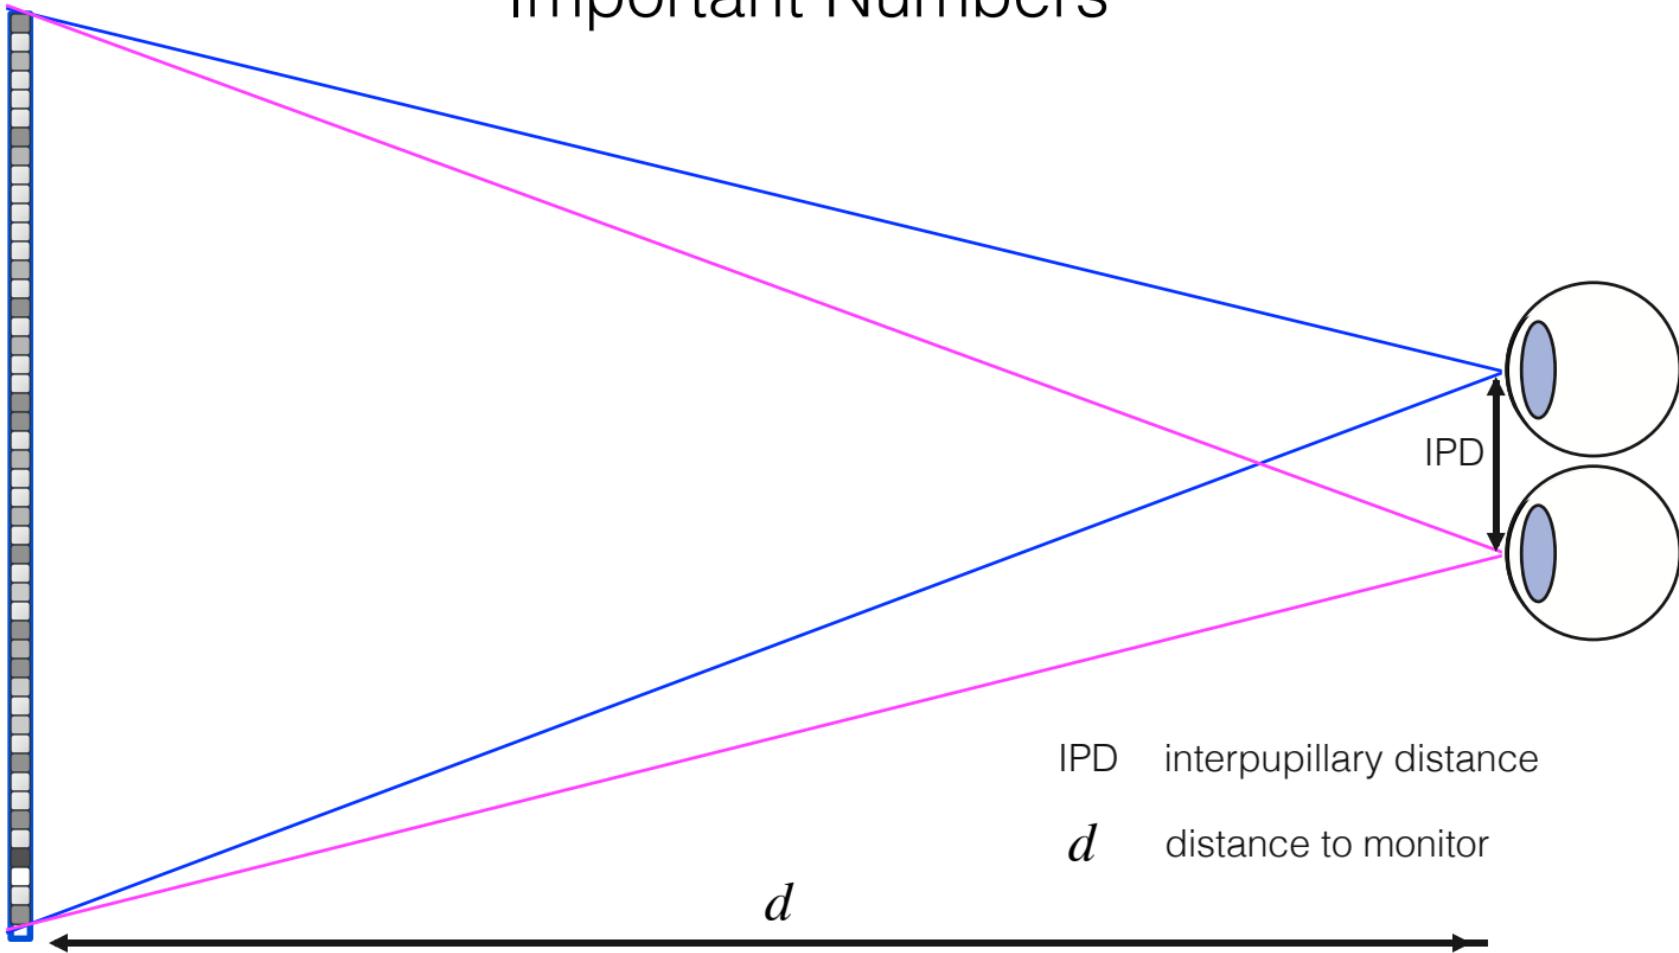
Take Off Your 3D Glasses Now!

# Stereo Rendering with OpenGL/WebGL: View Matrix

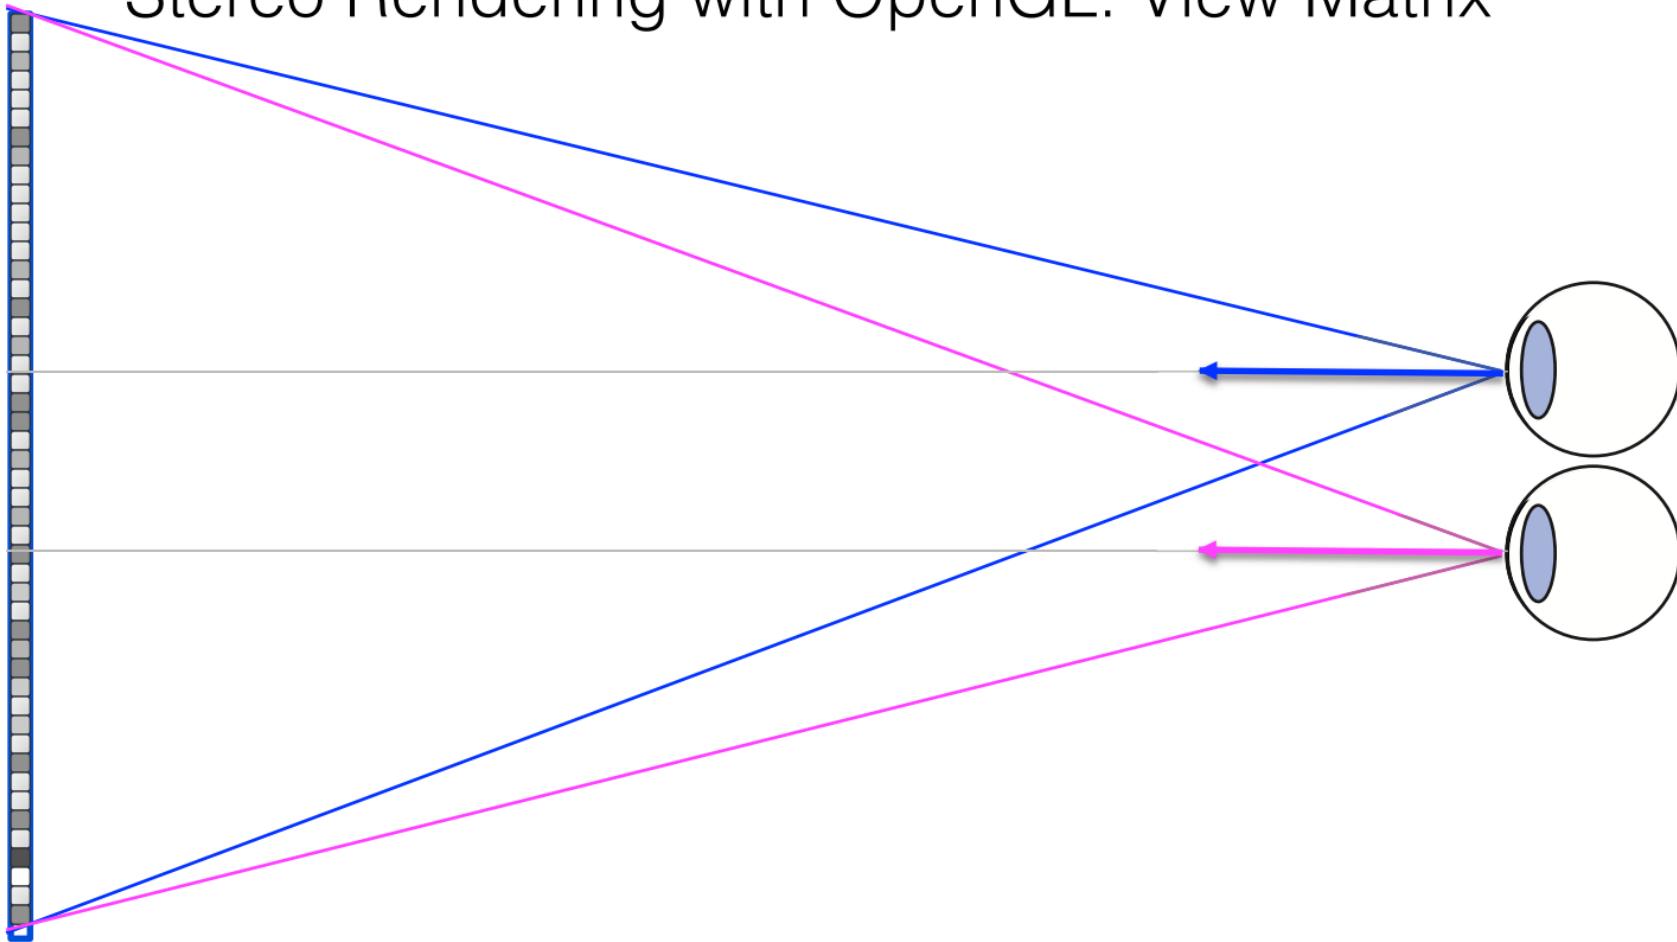
- need to modify view matrix and projection matrix
- rendering pipeline does not change – only those two matrices
- however: need to render two images in sequence (more details later)
- look at view matrix first: `THREE.Matrix4().lookAt(eye,center,up)`
- function is as useful as it has been, just need to adjust parameters

# Important Numbers

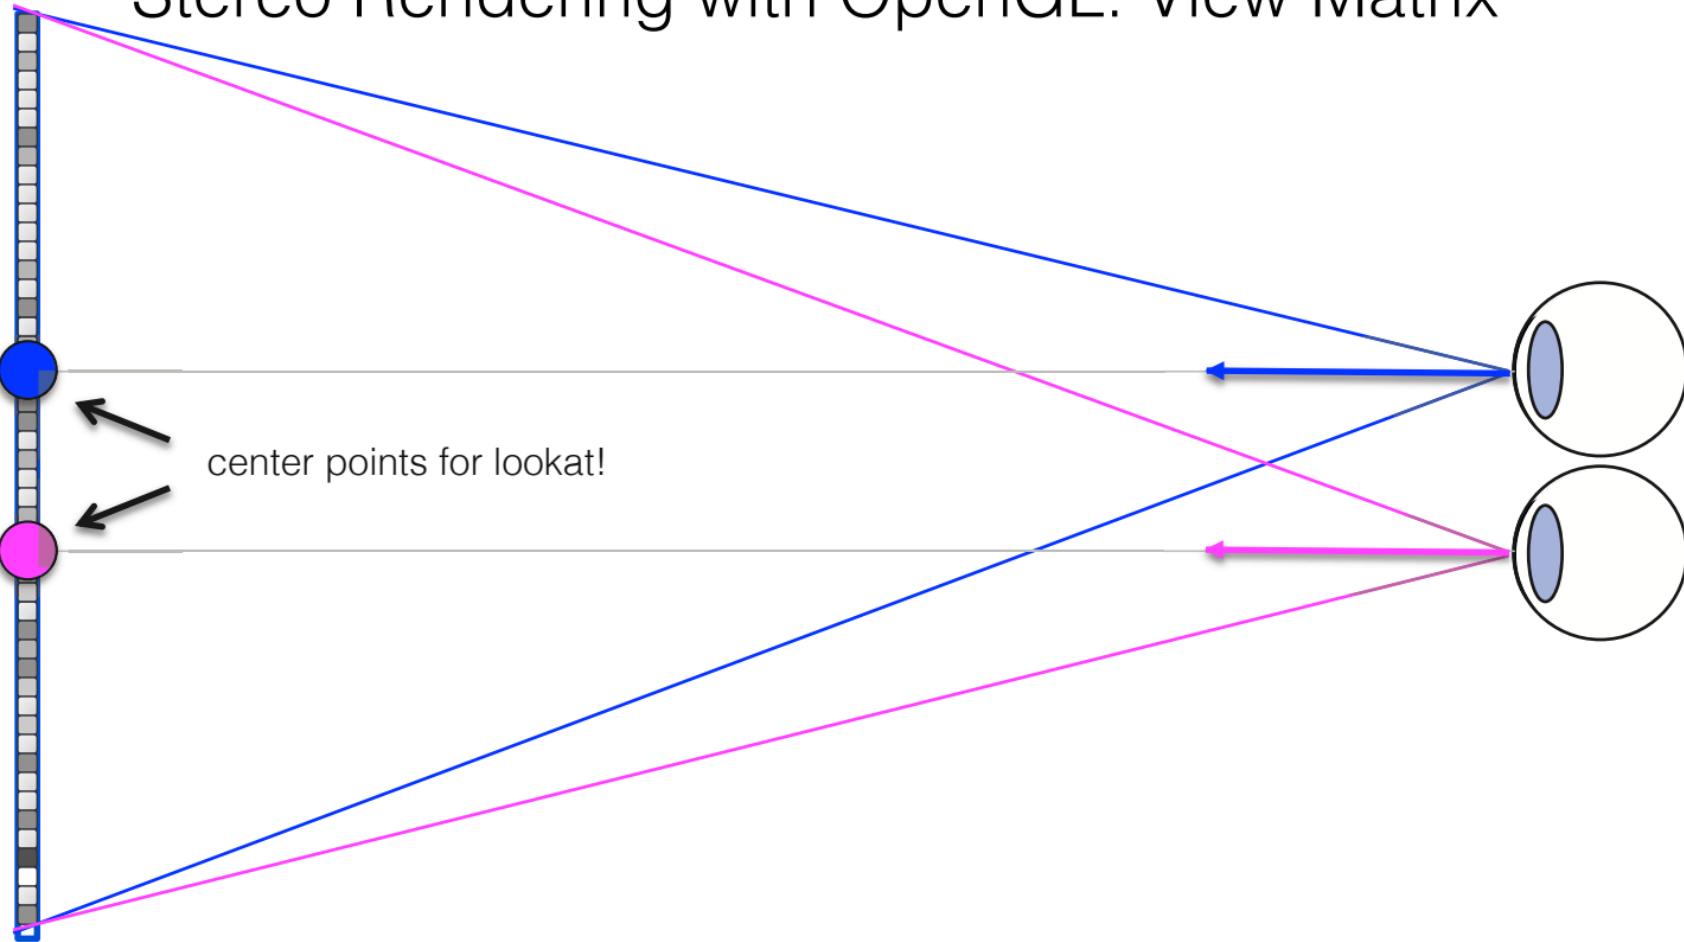
monitor width & height



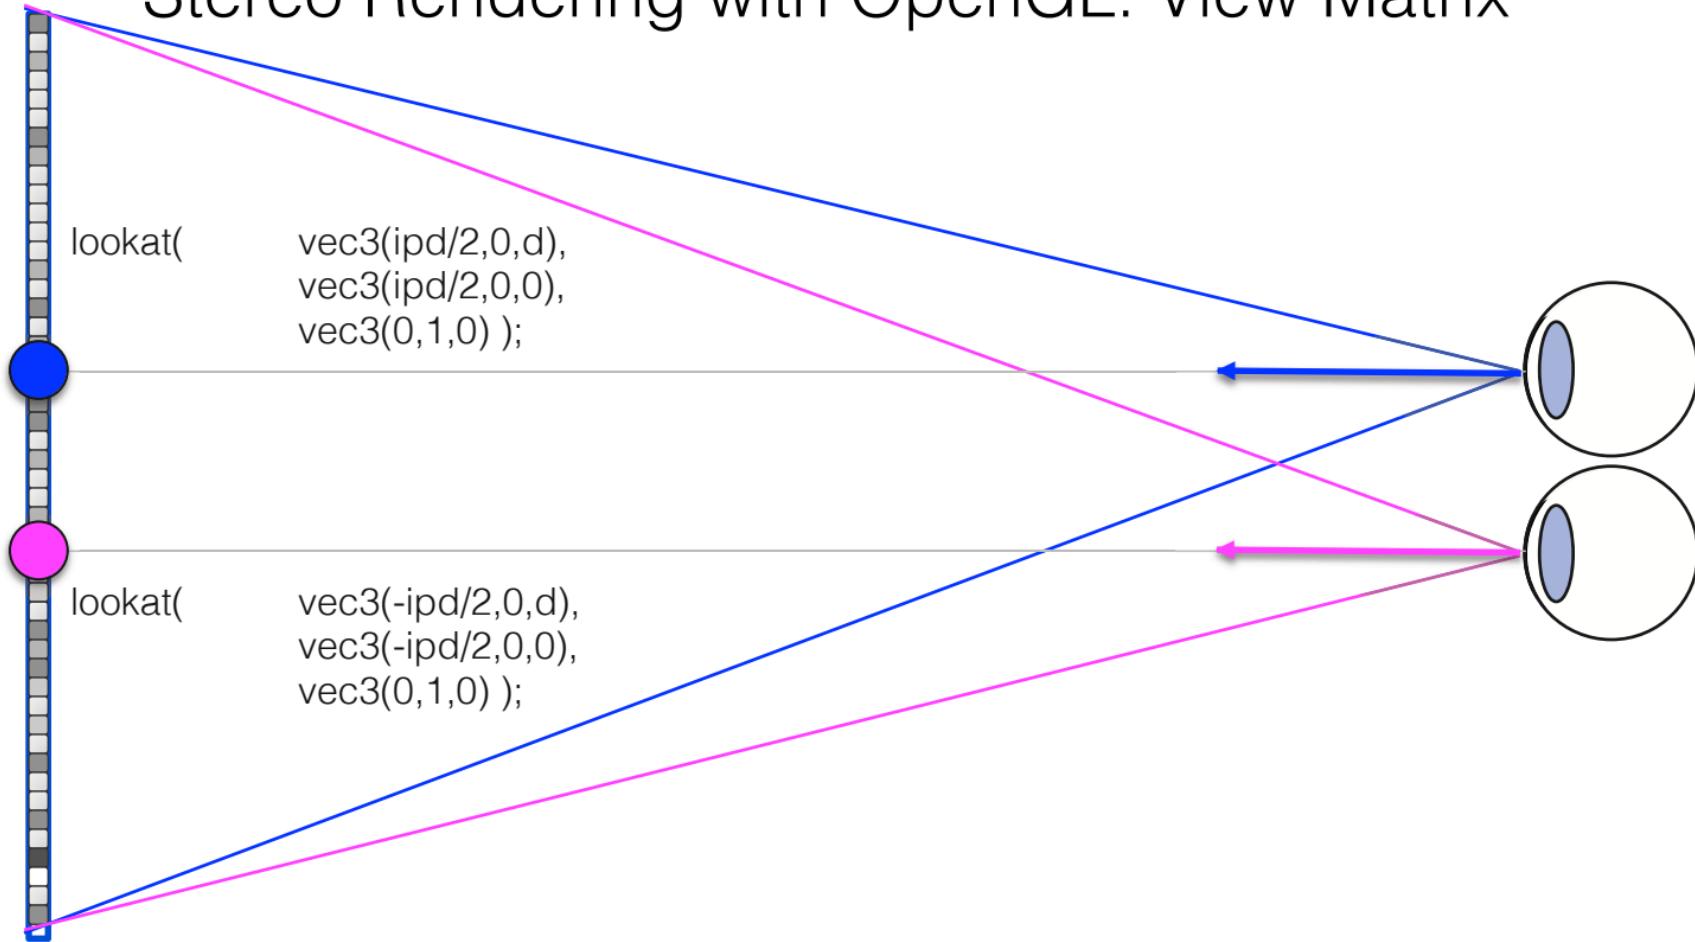
# Stereo Rendering with OpenGL: View Matrix



# Stereo Rendering with OpenGL: View Matrix

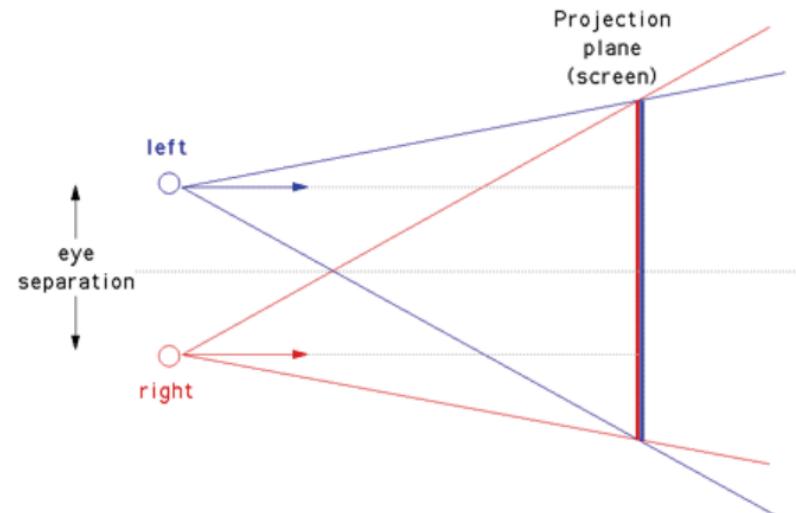


# Stereo Rendering with OpenGL: View Matrix

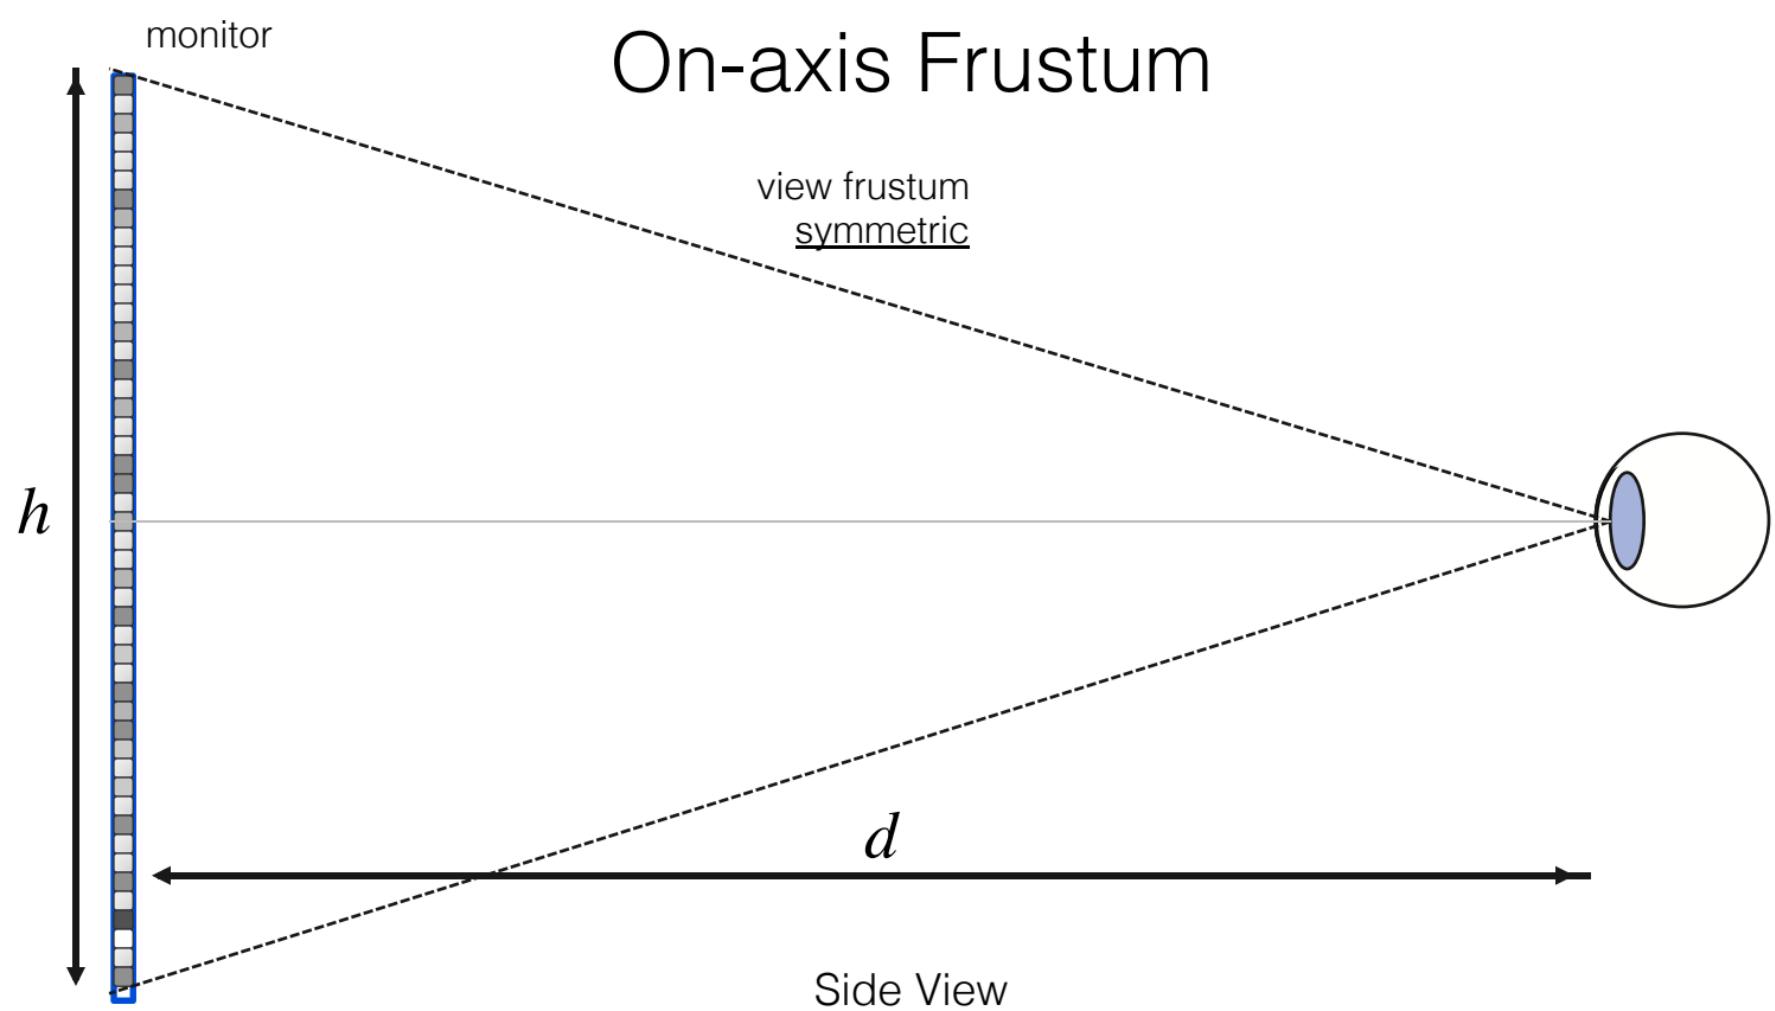


# Stereo Rendering with OpenGL: Projection Matrix

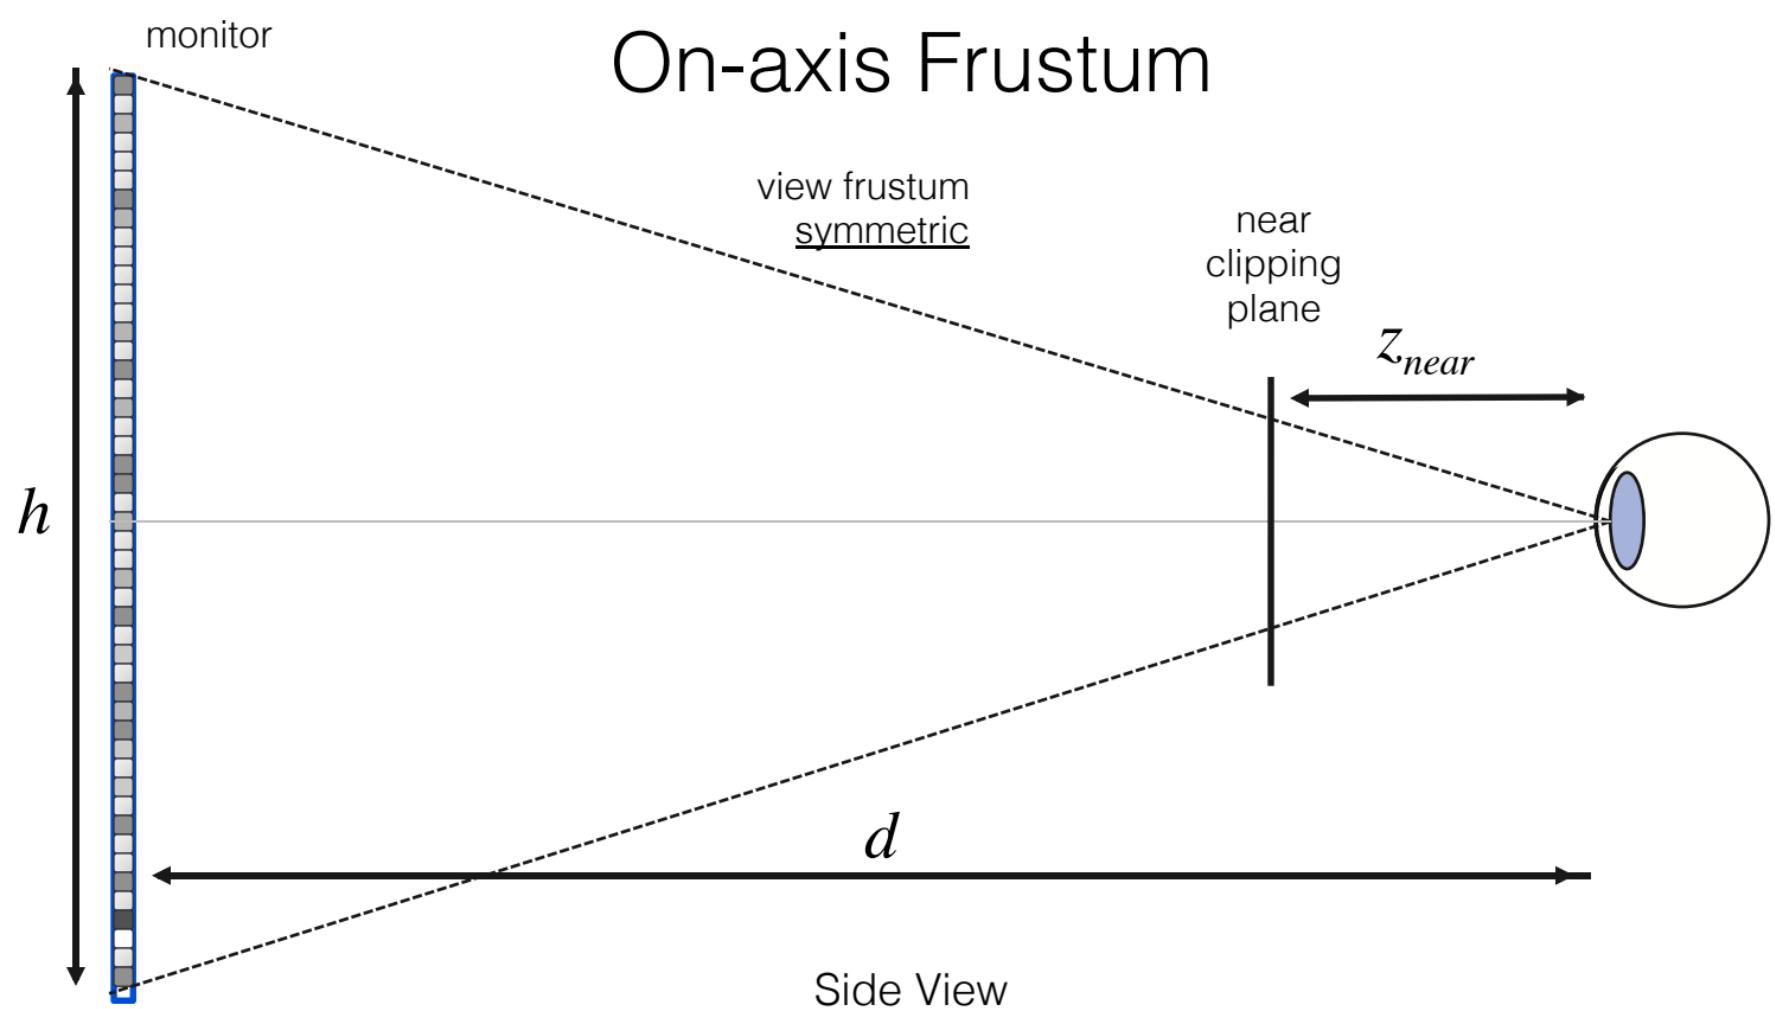
- perspective projection we have discussed so far is on-axis=symmetric
- we need a different way to set up the asymmetric, off-axis frustum
- use `THREE.Matrix4().makePerspective(left,right,top,bottom,znear,zfar)`



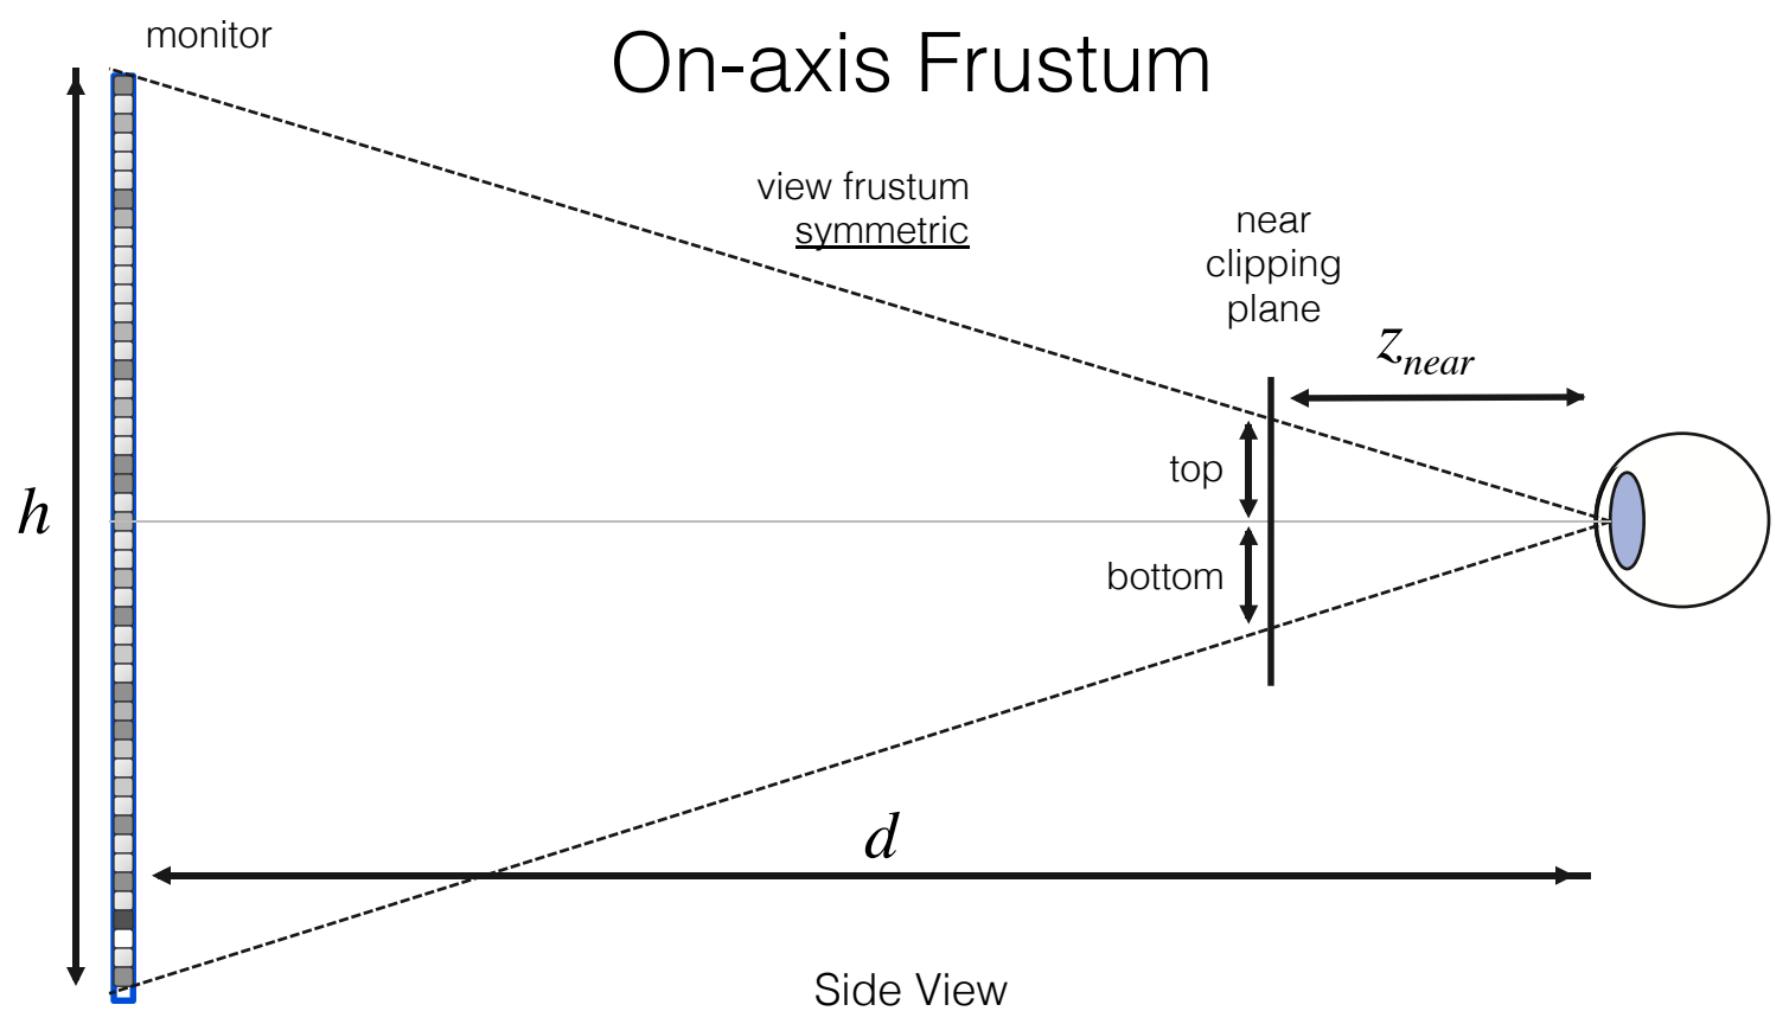
# On-axis Frustum



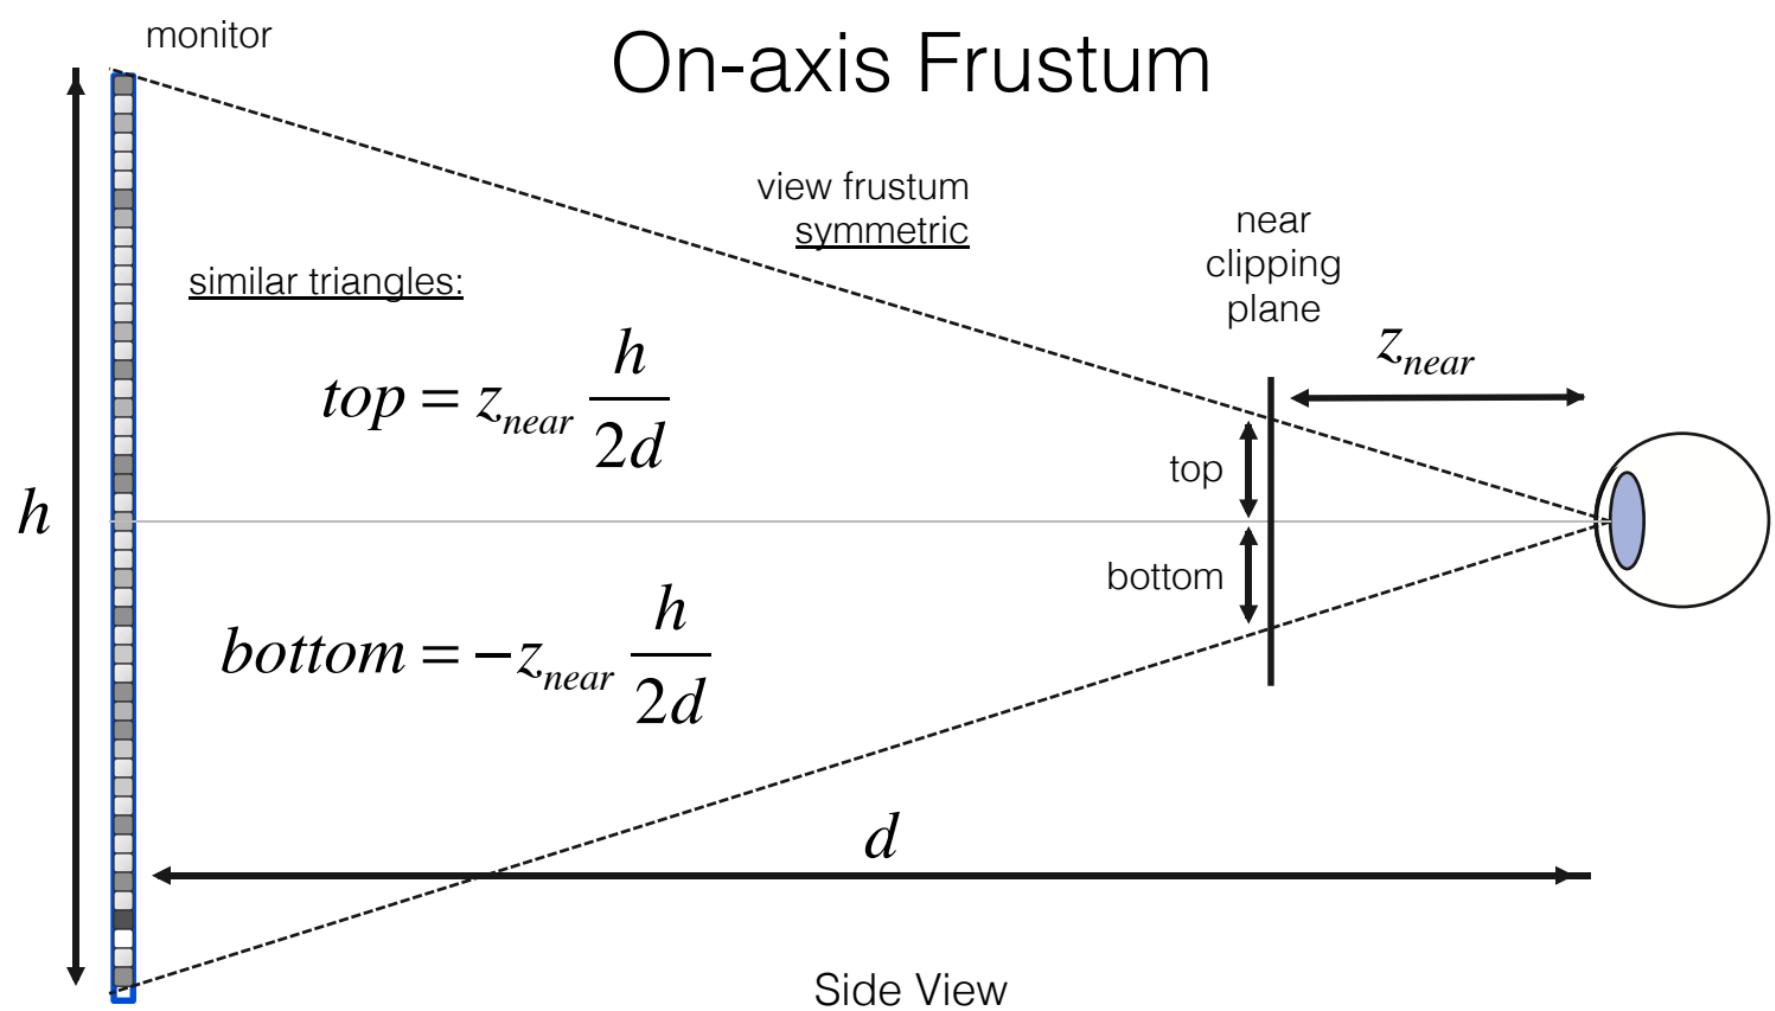
# On-axis Frustum



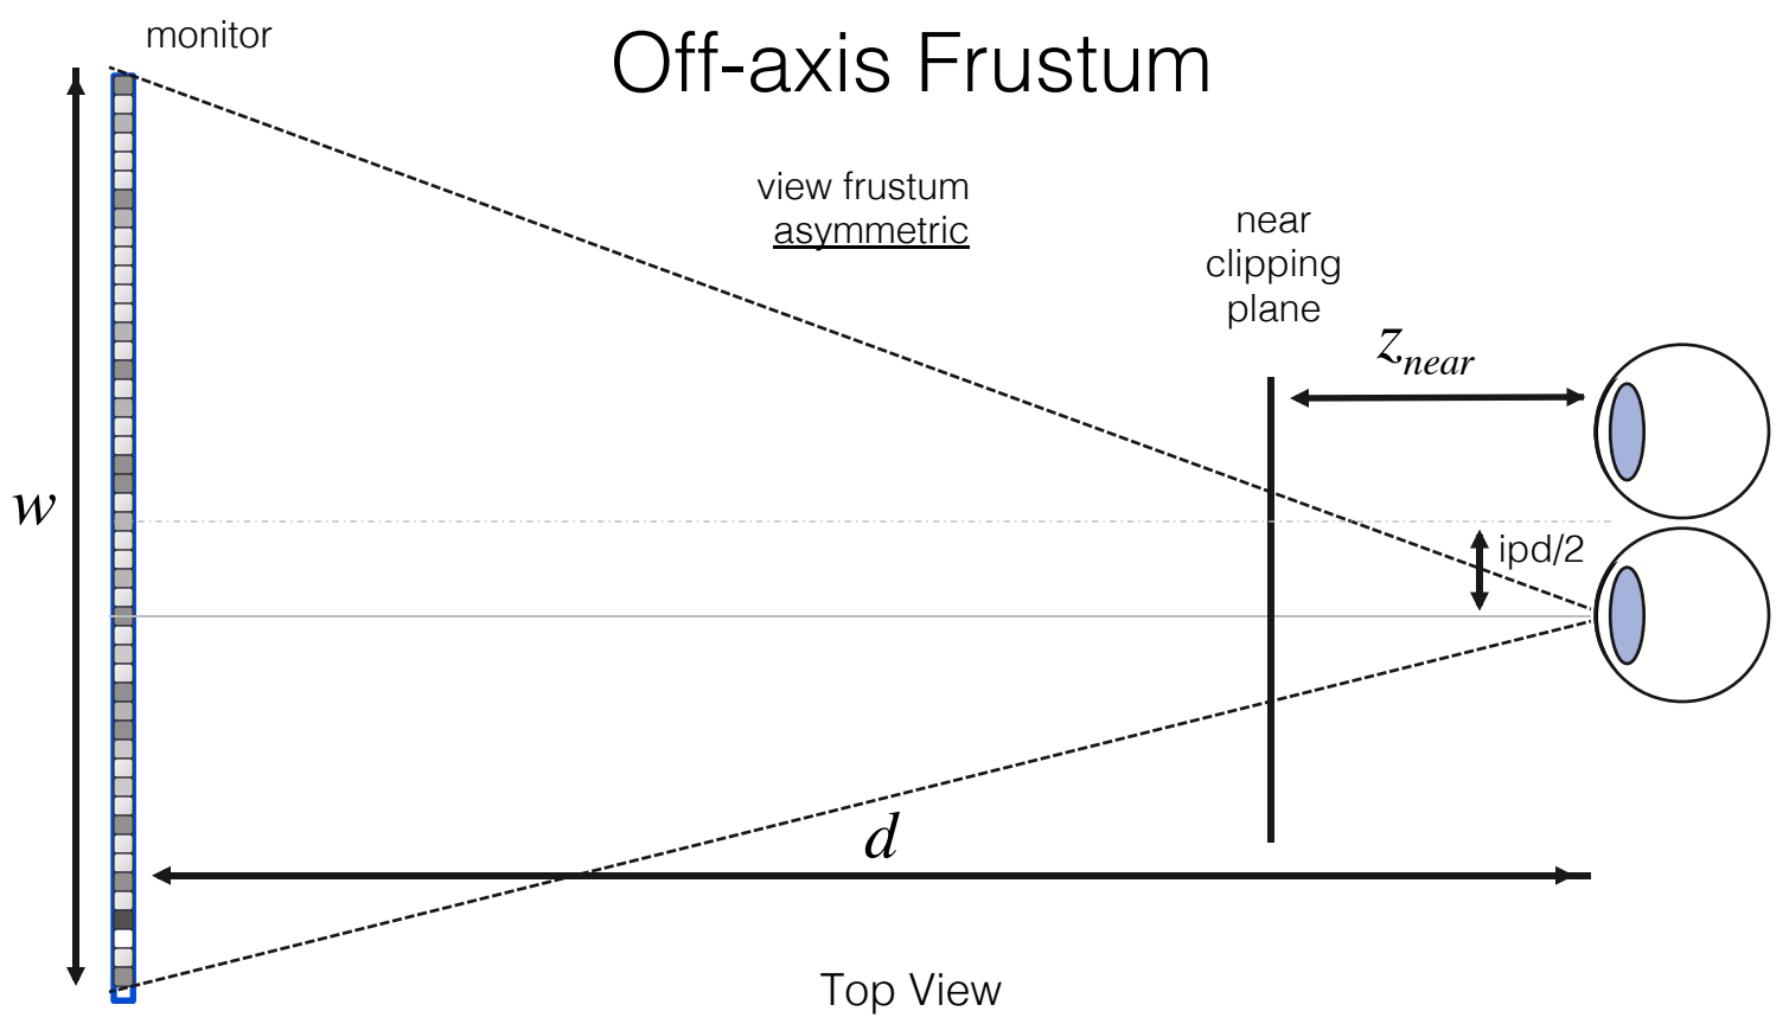
# On-axis Frustum



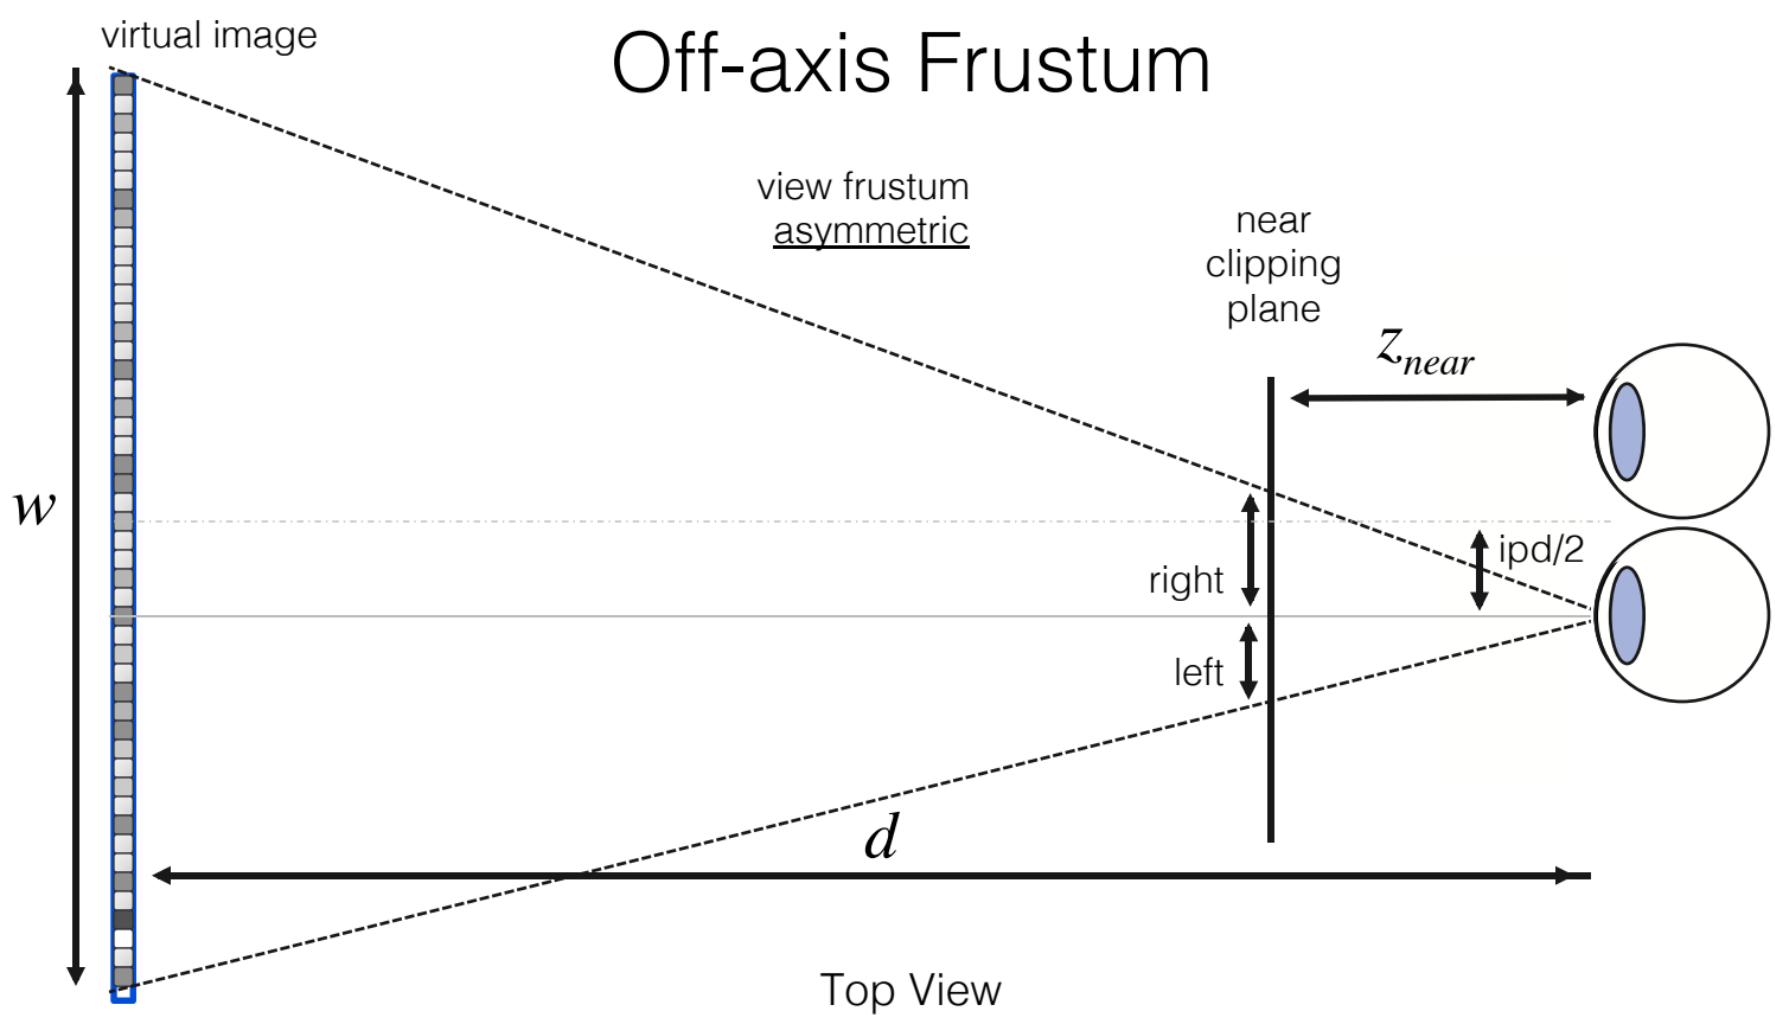
# On-axis Frustum



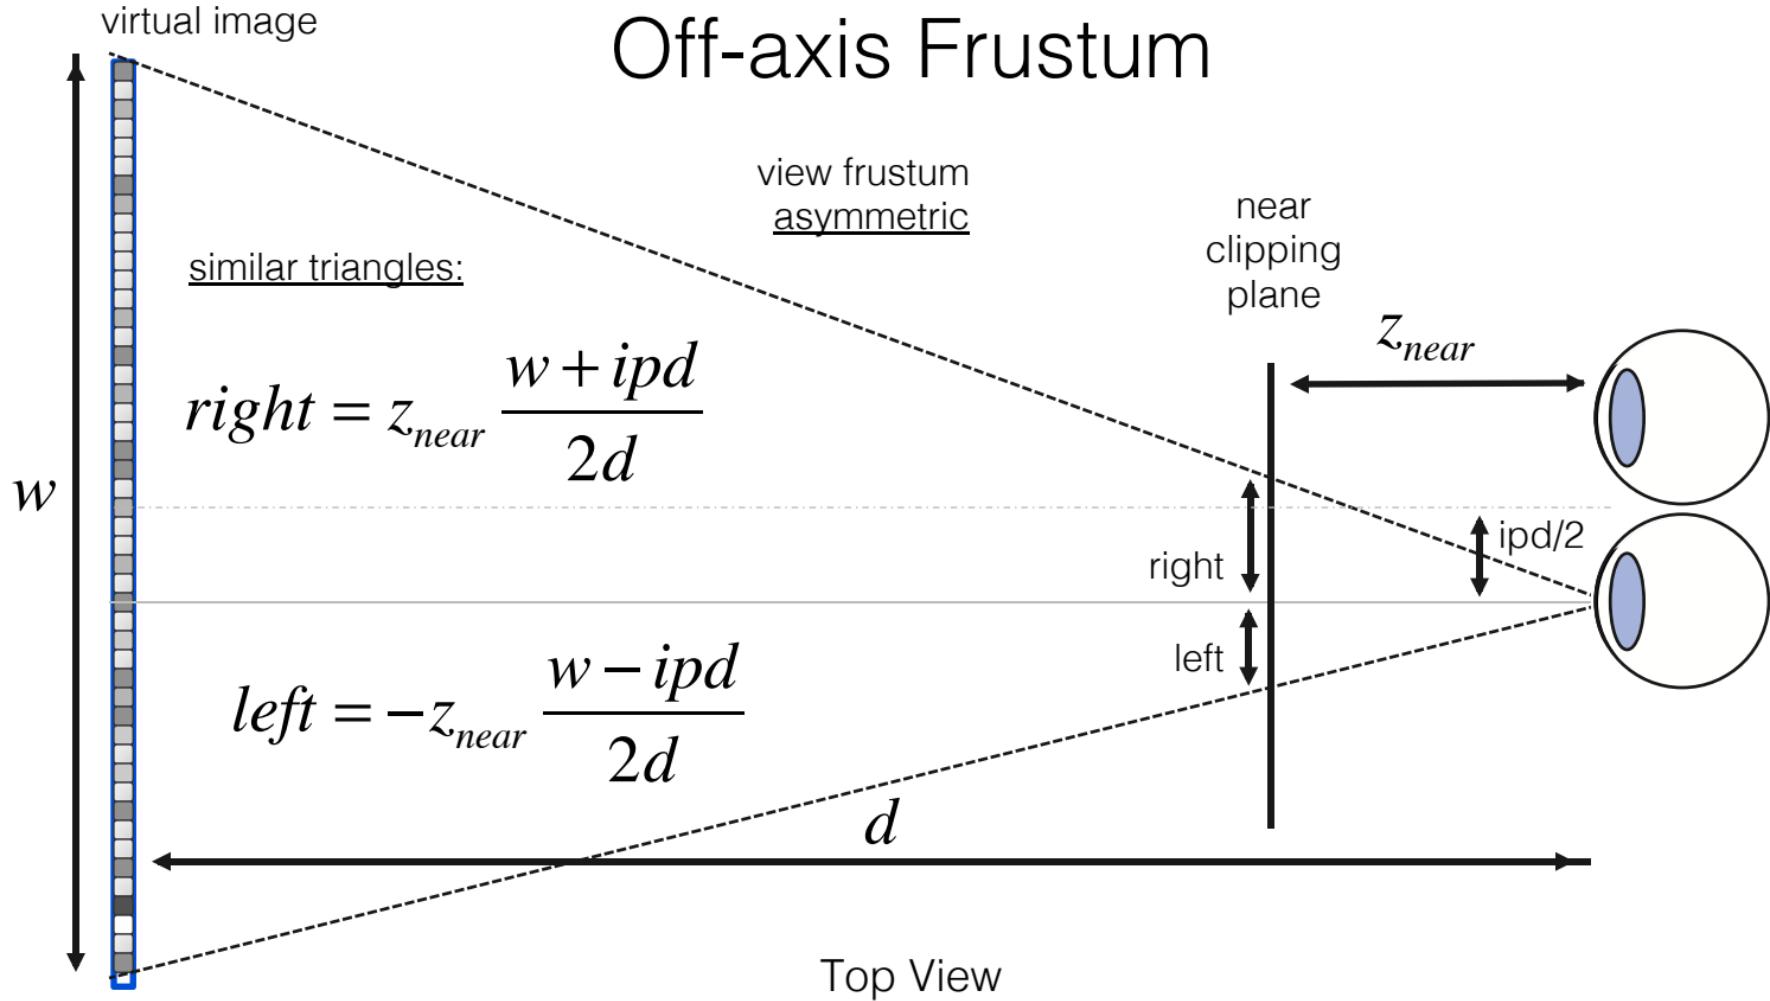
# Off-axis Frustum



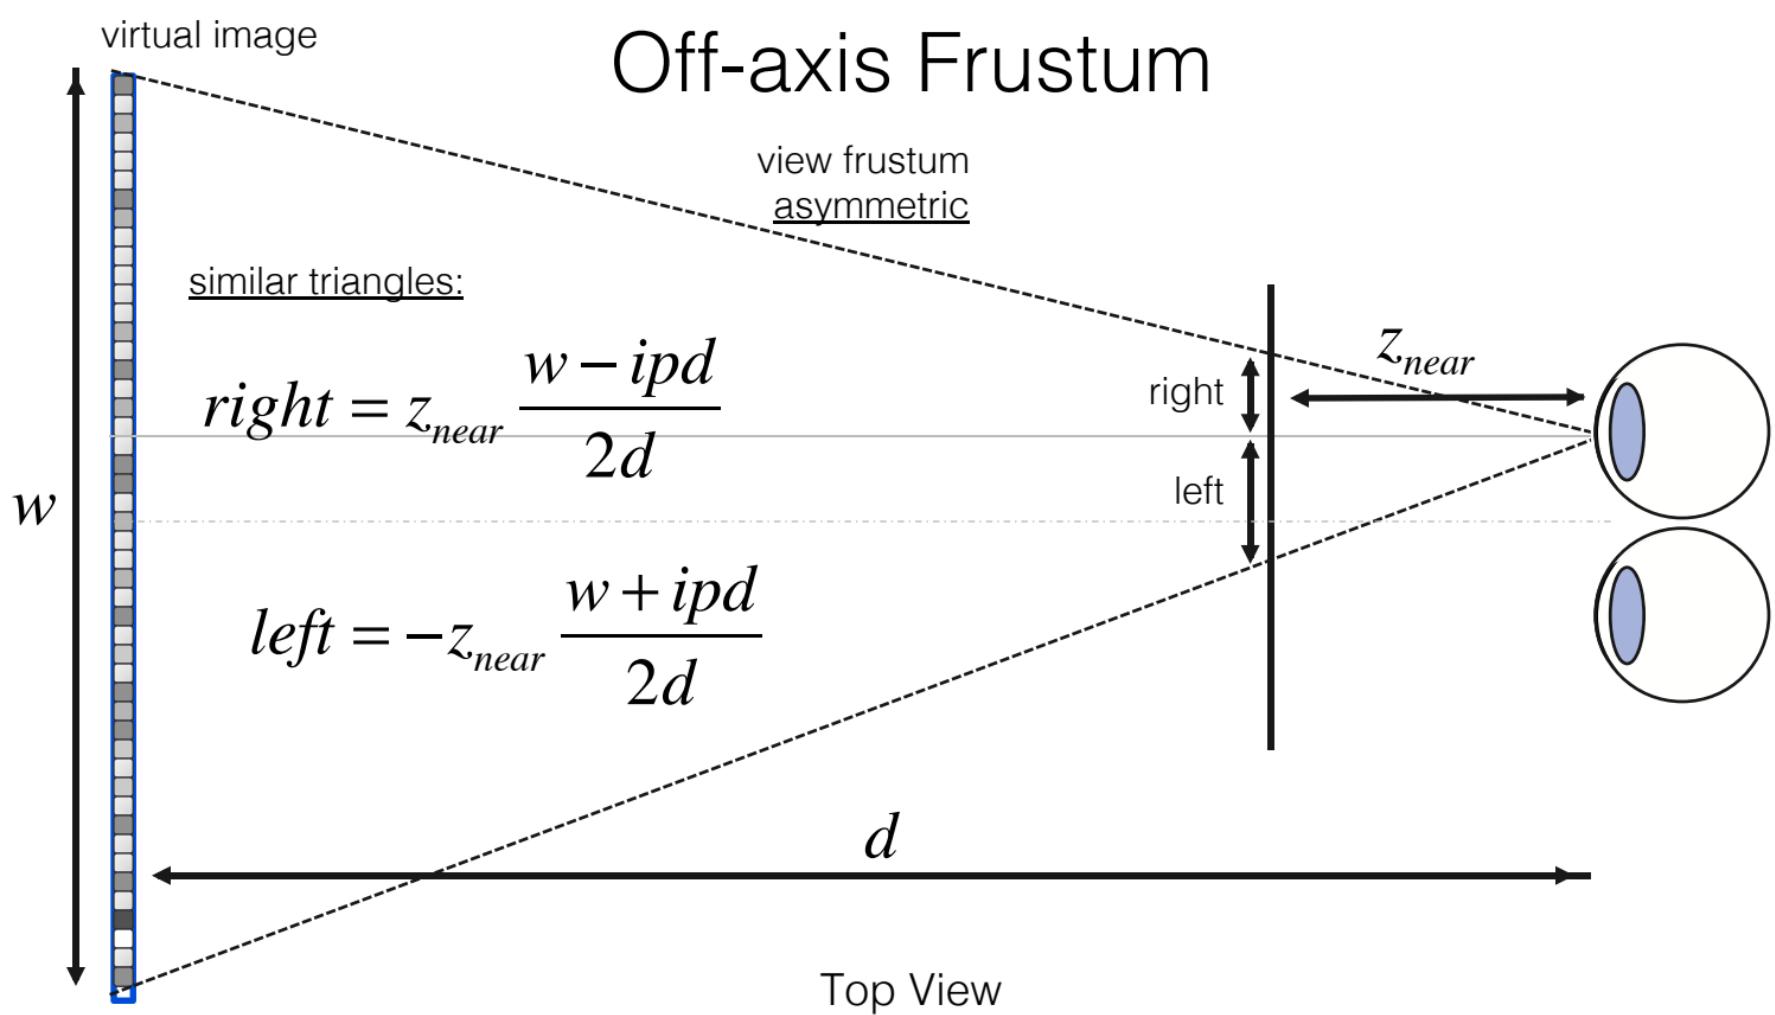
# Off-axis Frustum



# Off-axis Frustum



# Off-axis Frustum



# Anaglyph with OpenGL

- most efficient way:
  1. clear color and depth buffer
  2. set left modelview and project matrix, render scene only into red channel
  3. clear depth buffer
  4. set right modelview and project matrix, render scene only into green & blue channels
- we'll do it in a slightly more complicated way (need next week anyway):
  - multiple render passes
  - render into offscreen (frame) buffers

# OpenGL Frame Buffers

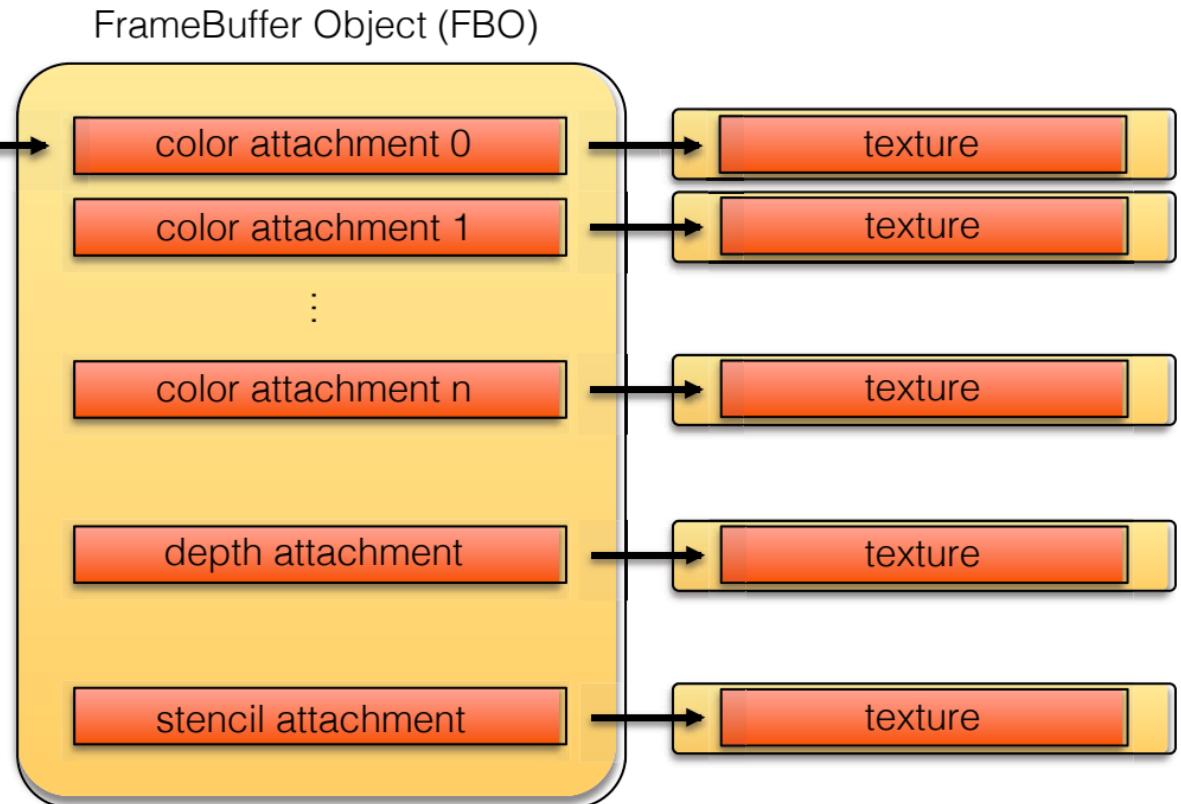
- usually (frame) buffers are provided by the window manager (i.e., your browser)
- for most mono applications, two (double) buffers: back buffer and front buffer
- render into back buffer; swap buffers when done rendering
- advantage: rendering takes time, you don't want the user to see how triangles get drawn
- in many stereo applications, 4 (quad) buffers: front/back left and right buffer
- render left and right images into back buffers, then swap both together

# OpenGL Frame Buffers

- more generic model: offscreen buffer
- most common form of offscreen buffer in OpenGL: framebuffer object
- concept of “render-to-texture” but with multiple “attachments” for color, depth, and other important per-fragment information
- as many framebuffer objects as desired, they all “live” on the GPU (no memory transfer)
- bit depth per color: 8 bits, 16 bits, 32 bits for color attachments; 24 bits for depth

# OpenGL Frame Buffers

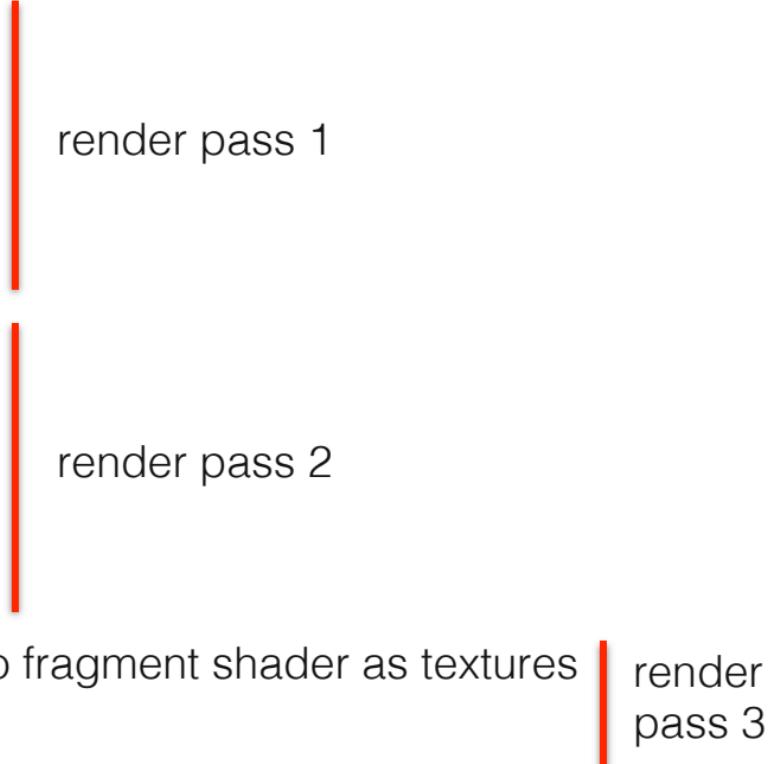
- render into FBO as usual, just enable/disable the FBO
- access content by texture ID (e.g. in GLSL shader)



# OpenGL Frame Buffers

- FBOs are crucial for multiple render passes!
- 1<sup>st</sup> pass: render color and depth into FBO
- 2<sup>nd</sup> pass: render textured rectangle – access FBO in fragment shader
- we'll provide a simple-to-use interface that shields you from the details of FBOs
- more details in lab on Friday ...

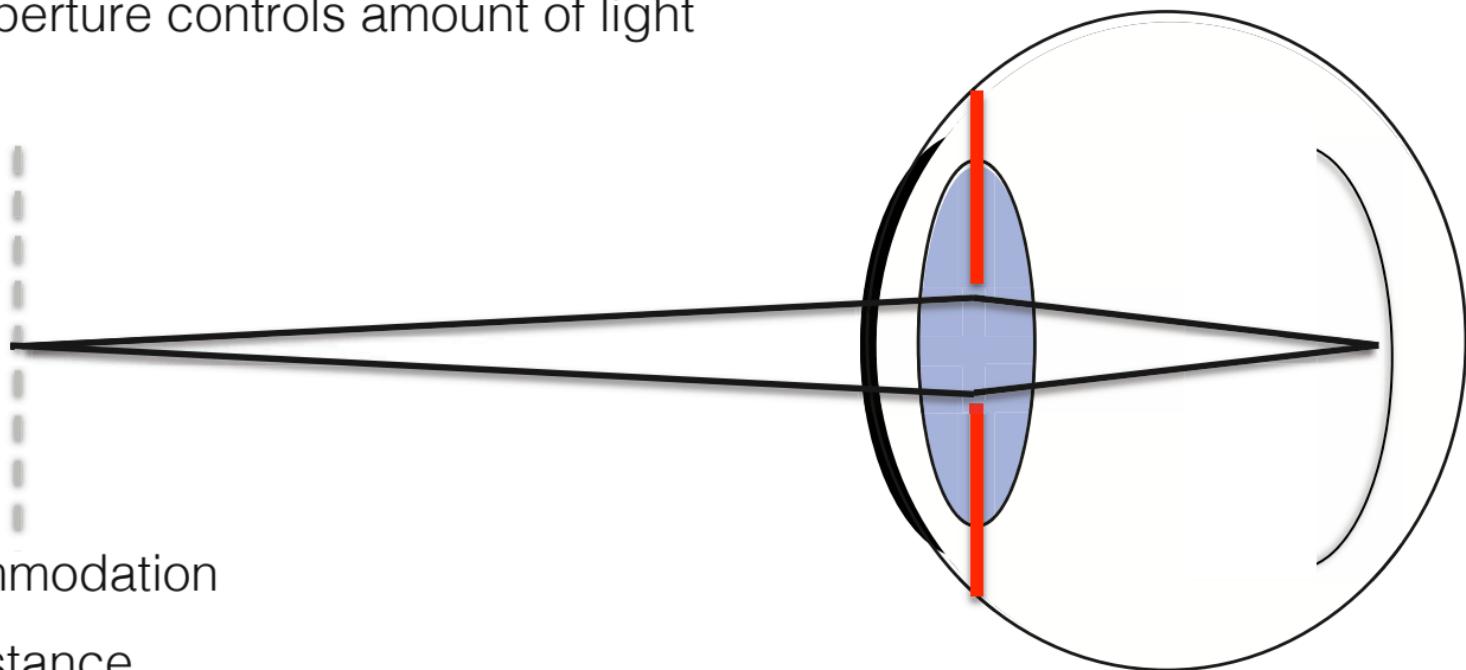
# Anaglyph Rendering with OpenGL & GLSL

1. activate FBO1
  2. set left modelview & projection matrix
  3. render scene
  4. deactivate FBO1
  5. activate FBO2
  6. set right modelview & projection matrix
  7. render scene
  8. deactivate FBO2
  9. render rectangle, pass FBO1 and FBO2 into fragment shader as textures
  10. merge stereo images in fragment shader
- 
- render pass 1
- render pass 2
- render pass 3

# Retinal Blur / Depth of Field Rendering

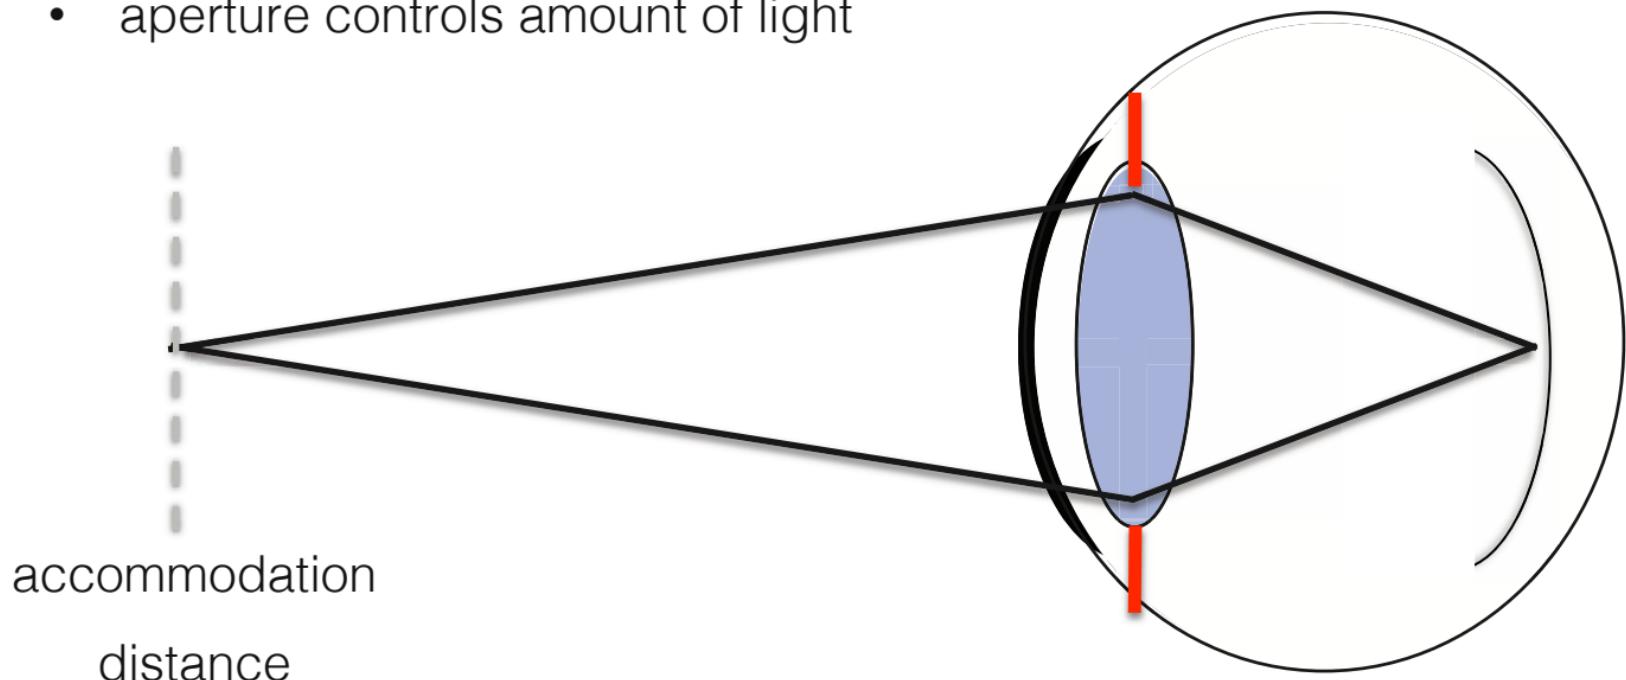
- aperture controls amount of light

accommodation  
distance



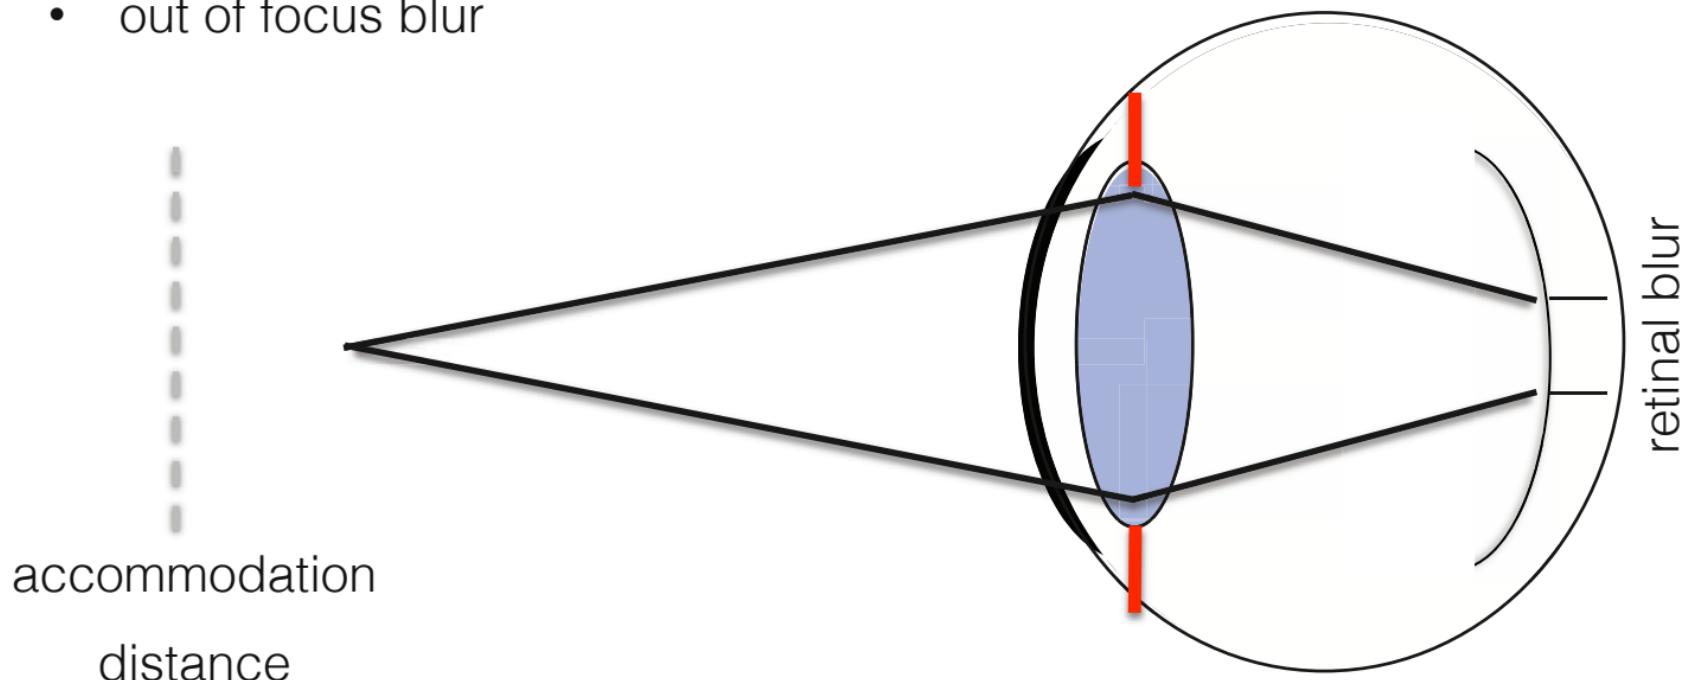
# Retinal Blur / Depth of Field Rendering

- aperture controls amount of light



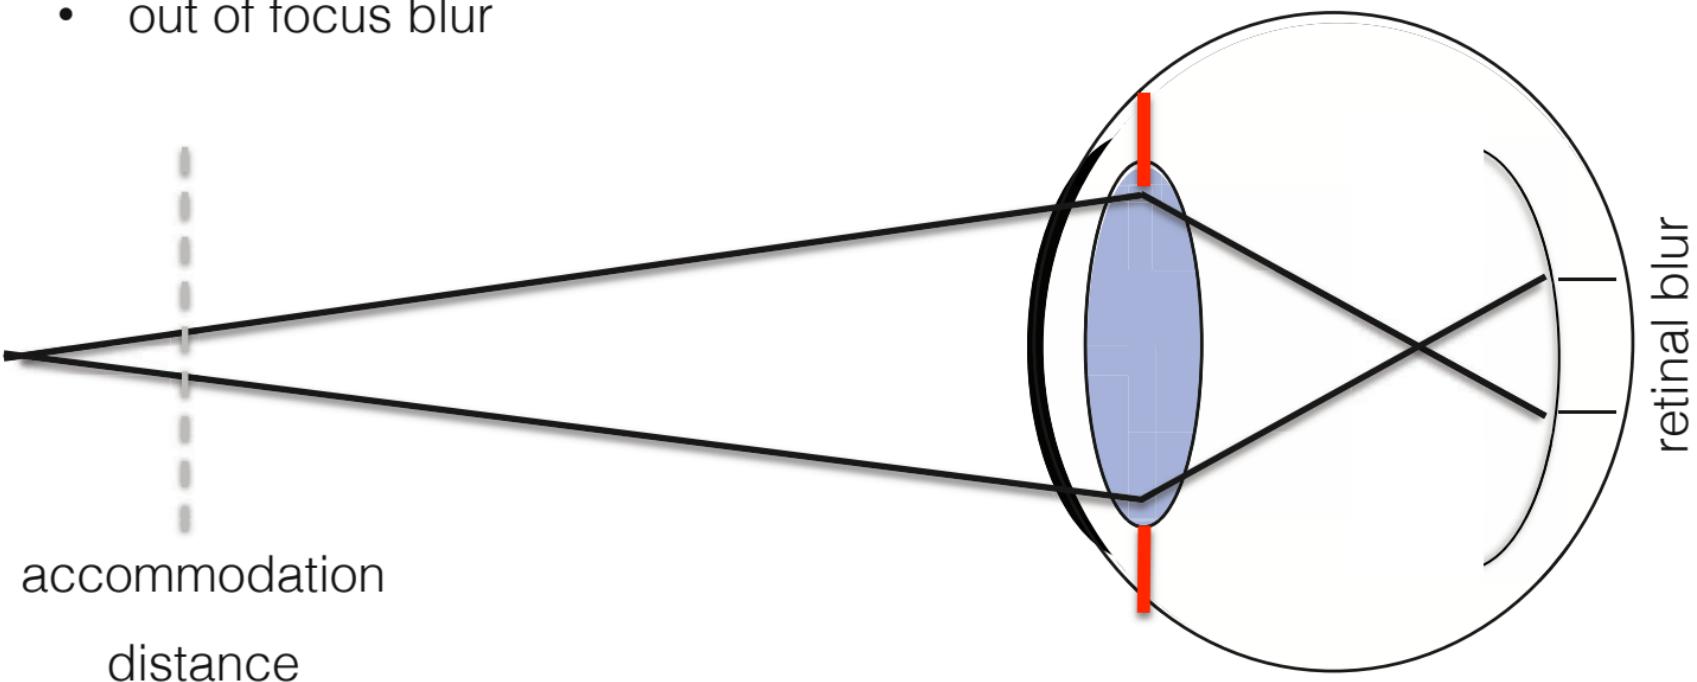
# Retinal Blur / Depth of Field Rendering

- out of focus blur



# Retinal Blur / Depth of Field Rendering

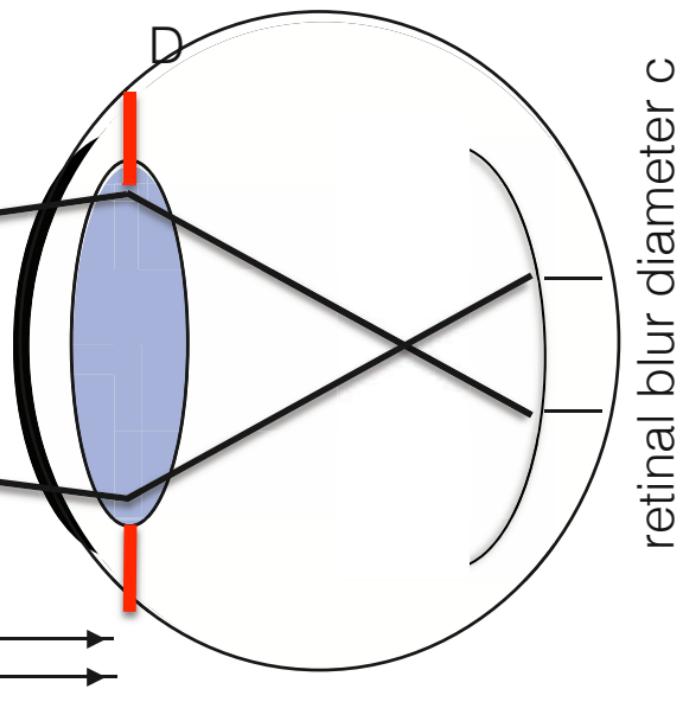
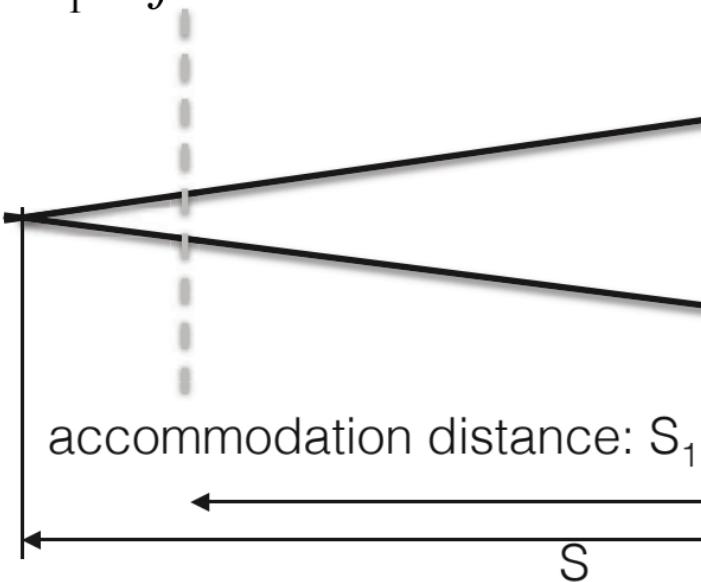
- out of focus blur



# Retinal Blur / Depth of Field Rendering

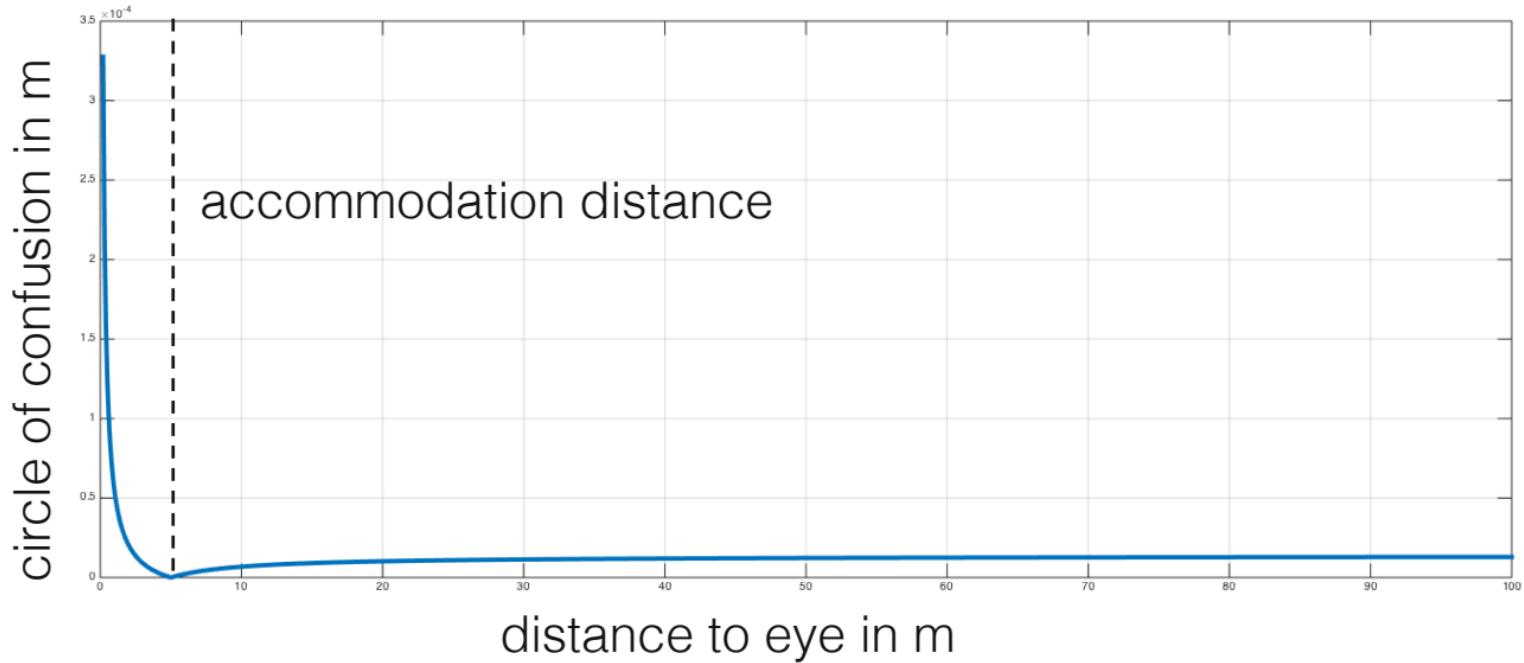
$$c = M \cdot D \cdot \frac{|S - S_1|}{S}$$

$$M = \frac{f}{S_1 - f} \quad f = 17\text{mm}$$



# Circle of Confusion

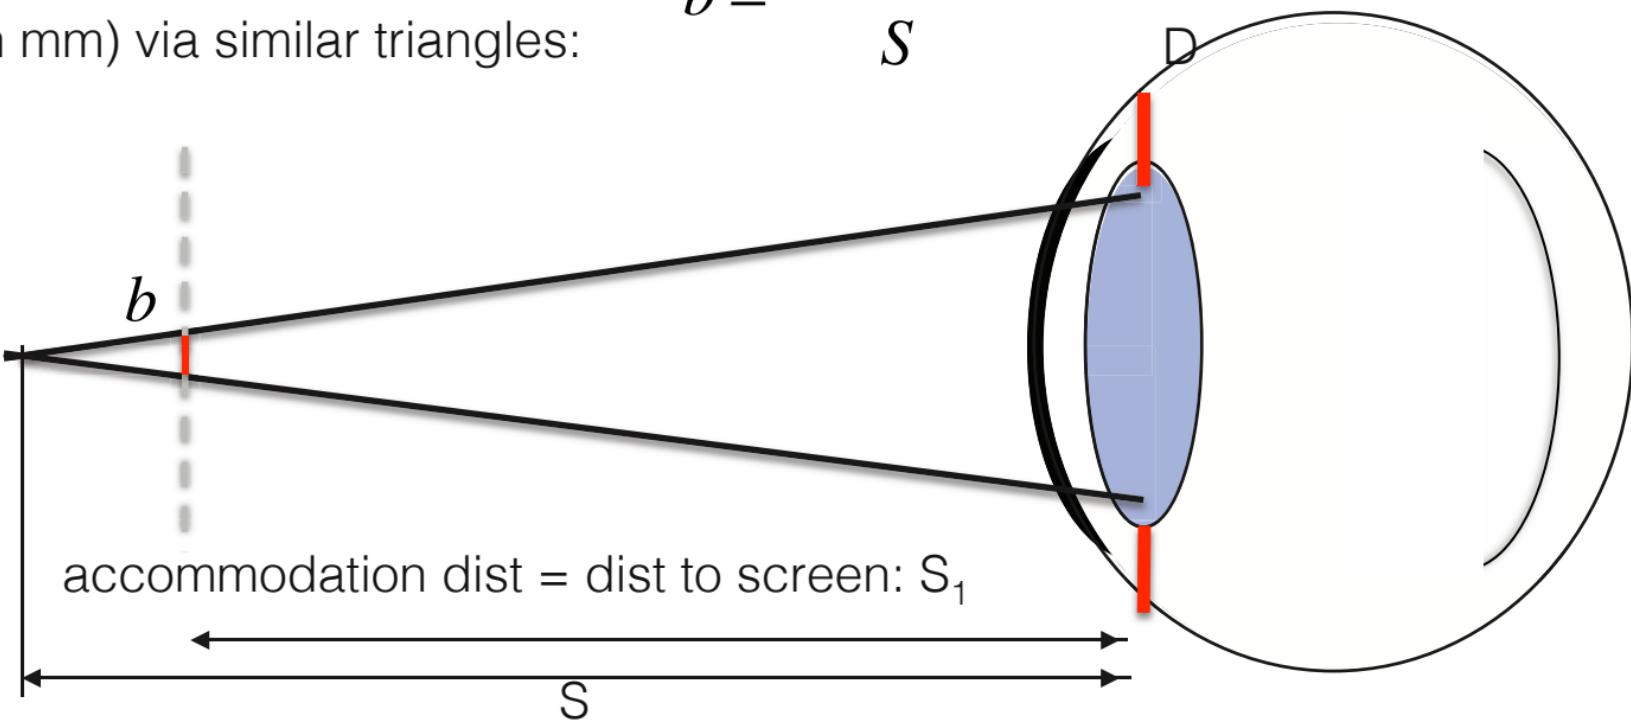
$$c = M \cdot D \cdot \frac{|S - S_1|}{S}$$



# Retinal Blur / Depth of Field Rendering

depth of field blur on screen  
(in mm) via similar triangles:

$$b = \frac{|S - S_1|D}{S}$$



# Depth of Field with OpenGL/GLSL

- two rendering passes:
  1. render image and depth map into FBO
  2. render quad textured with image + depth
    - vertex shader is pass-through (just transforms, pass on texture coordinates, no lighting)
    - in fragment shader:
      - calculate depth for each fragment in mm (not clip coords)
      - calculate retinal blur size in pixels given depth & pupil diameter
      - apply blur via convolution with double for loop over neighboring color values in the texture

# Depth of Field with OpenGL/GLSL

- how to get metric depth of a fragment ?
- in fragment shader we have `gl_FragCoord` with `gl_FragCoord.z` in [0,1]
- get normalized device coordinates, i.e.  $z_{NDC} = 2 * \text{gl_FragCoord.z} - 1$  in [-1,1]
- with multiple rendering passes, `gl_FragCoord.z` will be in depth map (output of 1<sup>st</sup> rendering pass, later accessed as texture)
- now undo perspective projection to get z in view coordinates:

$$\begin{pmatrix} x_{NDC} \\ y_{NDC} \\ z_{NDC} \\ 1 \end{pmatrix} = \begin{pmatrix} x_{clip} / w_{clip} \\ y_{clip} / w_{clip} \\ z_{clip} / w_{clip} \\ 1 \end{pmatrix}$$

$$\begin{pmatrix} x_{clip} \\ y_{clip} \\ z_{clip} \\ w_{clip} \end{pmatrix} = M_{proj} \cdot M_{view} \cdot M_{model} \cdot \underbrace{\begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}}_v = M_{proj} \cdot \underbrace{\begin{pmatrix} x_{view} \\ y_{view} \\ z_{view} \\ 1 \end{pmatrix}}_{v_{view}}$$

# Depth of Field with OpenGL/GLSL

- how to get metric depth of a fragment ?

$$M_{proj} = \begin{pmatrix} \frac{2n}{r-l} & 0 & \frac{r+l}{r-l} & 0 \\ 0 & \frac{2n}{t-b} & \frac{t+b}{t-b} & 0 \\ 0 & 0 & -\frac{f+n}{f-n} & \frac{-2 \cdot f \cdot n}{f-n} \\ 0 & 0 & -1 & 0 \end{pmatrix} \rightarrow z_{clip} = -\frac{f+n}{f-n} z_{view} - \frac{2fn}{f-n} \quad w_{clip} = -z_{view}$$
$$z_{view} = \frac{2fn}{f-n} \cdot \frac{1}{z_{NDC} - \frac{f+n}{f-n}}$$

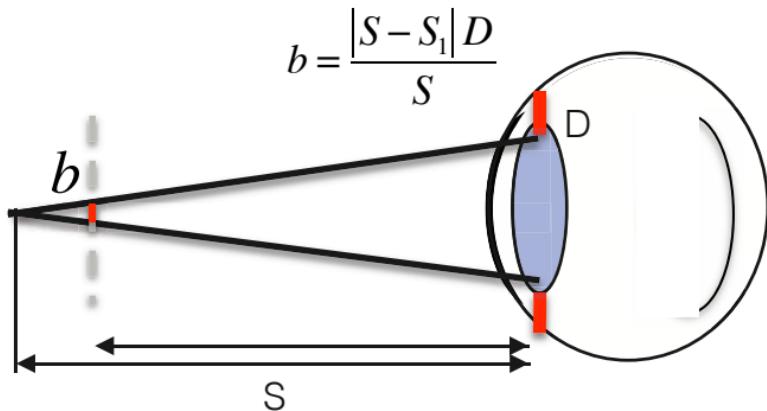
$$\underbrace{\begin{pmatrix} x_{NDC} \\ y_{NDC} \\ z_{NDC} \\ 1 \end{pmatrix}}_{v_{NDC}} = \begin{pmatrix} x_{clip} / w_{clip} \\ y_{clip} / w_{clip} \\ z_{clip} / w_{clip} \\ 1 \end{pmatrix} = M_{proj} \cdot M_{view} \cdot M_{model} \cdot \underbrace{\begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}}_v = M_{proj} \cdot \underbrace{\begin{pmatrix} x_{view} \\ y_{view} \\ z_{view} \\ 1 \end{pmatrix}}_{v_{view}}$$

# Depth of Field with OpenGL/GLSL

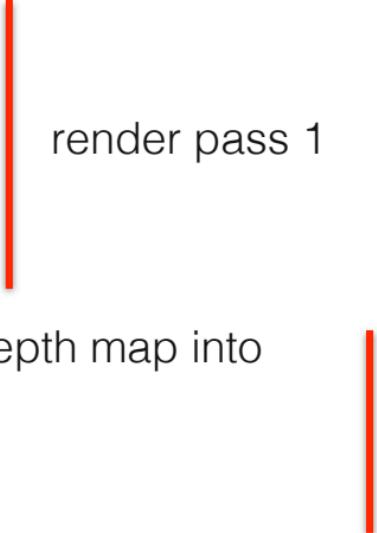
- how to compute retinal blur size and convert to pixels ?

$$pixel\_size_{x/y} = \frac{screen\_size_{x/y}}{screen\_resolution_{x/y}}$$

$$blur\_diameter\_px = \frac{b}{pixel\_size}$$



# Depth of Field with OpenGL/GLSL

1. activate FBO
  2. set modelview & projection matrix
  3. render 3D scene
  4. deactivate FBO
  5. render rectangle, pass FBO with image & depth map into  
fragment shader as textures
  6. execute depth of field fragment shader
- 
- render pass 1
- render pass 2

# Depth of Field with OpenGL/GLSL

- putting it all together – this is just a general overview, do not use this exact code

```
uniform sampler2D image;      // this was written in the first rendering pass as gl_FragColor
uniform sampler2D depthMap; // this was written in the first rendering pass as glFragDepth
uniform float      znear;
uniform float      zfar;
uniform float      pupilDiameter;
varying vec2       textureCoords;

void main () // fragment shader
{
    // get fragment z in NDC
    float zNDC = 2*sampler2D( depthMap, textureCoords ).r - 1;

    // get z in view coordinates (metric depth of current fragment)
    float zView = ...

    // compute retinal blur radius in pixels
    float blurRadius = ...
    int blurRadiusInt = round(blurRadius);

    // set output color by averaging neighboring pixels in the color image (i.e., convolution)
    gl_FragColor.rgb = 0;
    for (int i=-blurRadiusInt; i<blurRadiusInt; i++)
        for (int j=-blurRadiusInt; j<blurRadiusInt; j++)
            if (float(i*i+j*j) <= blurRadius*blurRadius)
                gl_FragColor.rgb += ... texture lookup in neighboring pixels

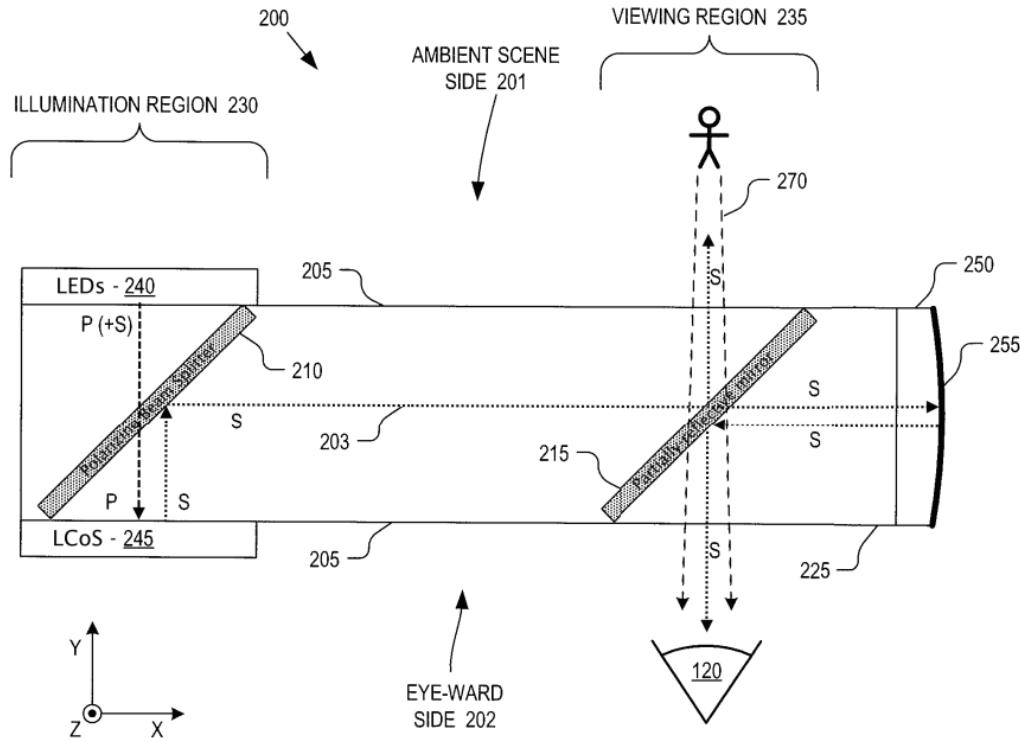
    // normalize color
    ...
}
```

# Summary

- many different technologies for glasses-based stereo
- we'll work with anaglyph for this lab + homework
- color management is important for anaglyph
- getting the view and projection matrices right is important (otherwise headaches)
- may need multiple render passes (more details in the lab)
- depth of field rendering may add more realism

# Next Lecture: HMD Optics and Microdisplays

- magnifiers
- VR & AR optics
- microdisplays
- stereo rendering for HMDs
- lens distortion / undistortion



drawing from Google Glass patent

# Further Reading

- <http://paulbourke.net/stereographics/stereorender/>
- Eric Dubois, “A Projection Method to Generate Anaglyph Stereo Images”, ICASSP 2001
- Library of Congress, Stereoscopic Cards:  
<http://www.loc.gov/pictures/search/?st=grid&co=stereo>